

DISTRIBUTED SYSTEMS

- Python
- Apache Beam
- Dataflow
- Pub/Sub
- Real-Time Streaming
- Low-Latency Services
- Autoscaling
- Traffic Routing

CLOUD & DEPLOYMENT

- GCP
- GKE
- Vertex AI
- Kubernetes
- Docker
- Helm
- Istio
- Seldon Core
- Azure ML
- Azure Functions

RELIABILITY & DELIVERY

- HPA
- Cloud Build
- ArgoCD
- Google KCC
- GitOps
- CI/CD
- Prometheus
- Grafana
- Load Testing
- Blue-Green Delivery

INFERENCE & GPU PLATFORMS

- vLLM
- LLM Serving
- LoRA / QLoRA
- PEFT
- Transformers
- Model Evaluation
- RAPIDS
- Distributed Training

CERTIFICATIONS

- Oracle Cloud Infrastructure 2024  
Generative AI Certified Professional
- Microsoft Certified  
Azure AI Engineer Associate
- Microsoft Certified  
Azure Data Scientist Associate
- Google Cloud  
Certified Generative AI Leader

EDUCATION

B.E. Electrical & Electronics  
Andhra University | 2012 - 2016

ML METHODS

- Scikit-learn
- PyTorch
- TensorFlow
- Keras
- DenseNet
- Classification
- Predictive Modelling
- Unsupervised Learning
- Feature Engineering
- Transfer Learning
- Anomaly Detection
- NLP
- Time-Series Forecasting

Production AI infrastructure engineer with 10+ years building distributed systems under real traffic. At Best Buy, architected shared LLM fine-tuning, evaluation, and vLLM serving on GKE with autoscaling, Istio routing, and blue-green delivery, with estimated annual API-cost savings of ~\$300K; also owns a streaming ML platform with feature pipelines handling 40,000 events/sec at peak and a 200 ms scoring SLA. Earlier led billion-row GPU pipelines on 8xA100 DGX. Hands-on with Python, Kubernetes, GCP, production inference, observability, performance diagnosis, and cross-team technical leadership.

EXPERIENCE

Lead Artificial Intelligence Engineer

Best Buy India  
Jul 2025 - Present | Bengaluru

- Architected a shared Python platform for config-driven LoRA/QLoRA training, automated evaluation, artifact lifecycle, and vLLM OpenAI-compatible serving on GKE; adapters hot-load into a shared base-model deployment.
- Designed autoscaled inference delivery with HPA, Istio traffic routing, blue-green releases, Cloud Build, ArgoCD, Google KCC, Prometheus, and Grafana; self-hosting supports estimated annual API-cost savings of ~\$300K. Kept model inputs company-managed and reduced external API exposure during peak demand.
- Built fraud systems across batch, streaming, and real-time models. Beam/Dataflow feature pipelines handled 40,000 events/sec at peak; supported a 200 ms scoring SLA with secure, monitored GKE/Seldon Core inference and load testing.
- Lead technical direction across data science, platform, fraud, and security teams; review designs and mentor engineers.

Lead ML Engineer — Machine Learning & MLOps

Wipro  
Aug 2023 - Jul 2025 | Hyderabad

- Built distributed training and inference pipelines on NVIDIA DGX (8xA100), processing billion-row AWS CloudTrail datasets with Morpheus, RAPIDS, continuous retraining, and Grafana alerting.
- Engineered phishing classification at 2,400 emails/sec - 2x prior throughput - and multifactor abuse-inbox processing integrated with Microsoft Sentinel and production security workflows.
- Led and mentored ML/MLOps engineers; coordinated client/NVIDIA delivery; completed 300+ résumé screens and conducted 50+ technical interviews.

Senior Data Scientist — Data Science / Azure ML

Accenture  
May 2021 - Sep 2023 | India

- Led end-to-end Azure ML product delivery for BamaGruppen, increasing inspections from 130 to 160 trucks/day, improving low-quality detection 17%, and generating \$2.5M annual savings.
- Built production services and automation with Python, Azure Functions, Logic Apps, Azure DevOps, CI/CD, and Google OR-Tools; translated ambiguous client needs into operated systems.

PROJECTS

LLM Training & Serving Platform  
Best Buy | GKE | vLLM | Autoscaling

Shared Python platform for training, evaluation, artifact management, model caching, adapter-aware serving, traffic routing, blue-green rollout, and full-stack observability.

Real-Time Retail ML Platform  
Best Buy | 40k feature events/sec | 200 ms scoring SLA

Streaming feature engineering and low-latency inference across high-traffic retail workflows, with autoscaling, load testing, secure endpoints, telemetry, and GitOps delivery.

GPU STACK

- CUDA GPU Libraries
- RAPIDS cuDF / CuPy
- Multi-GPU Training

OPTIMISATION & SIMULATION

- Google OR-Tools
- Vehicle Routing
- Mixed-Integer Programming
- Constraint Modelling
- Operations Research
- Assignment Models
- Supply-Chain Optimization
- Scenario Analysis

CLOUD & MANAGED ML

- GCS
- Artifact Registry
- Azure DevOps
- Function Apps
- Logic Apps

DATA & ML LIFECYCLE

- Vertex AI Feature Store
- Streaming Pipelines
- Training Pipelines
- Automated Evaluation
- Experiment Tracking
- Model Registry
- Model Versioning
- Continuous Training

LANGUAGES & DATA

- SQL
- Pandas
- NumPy
- MS SQL Server
- IBM DB2
- Power BI

ENGINEERING LEADERSHIP

- Architecture Reviews
- Technical Mentoring
- Technical Hiring
- Team Building
- Stakeholder Alignment
- Cross-Functional Delivery

ACHIEVEMENTS

15 honors and awards across Best Buy, Wipro, Accenture, and Infosys.  
Best Buy Talent Quotient (2026): cybersecurity and fraud delivery.  
Best Buy client AML Hall of Fame (2024): DGX delivery during the Wipro engagement.

ADDITIONAL PROJECTS & EXPERIENCE

Cybersecurity GPU Pipelines

Wipro | DGX | Billion-row scale

Continuously refreshed behavioral models and high-throughput classification pipelines with distributed GPU acceleration, monitoring, and production integrations.

vLLM Model Chooser

Open source | 115 models | A100-B200

Built a live GPU-to-model planner that filters vLLM-compatible models by VRAM, quantization, and KV-cache fit across accelerator classes.

EARLIER EXPERIENCE & ACHIEVEMENTS

Data Scientist

GD Research Centre | May 2020 - May 2021

Built domain-aware time-series forecasting and scenario analysis across energy, manufacturing, healthcare, finance, retail, telecommunications, automotive, and other industry intelligence centers.

Deep Learning Engineer

Tata Consultancy Services | Jul 2019 - Mar 2020

Built and deployed an end-to-end semiconductor ML system: distributed training on ~1 million images in ~3 hours, with 90% classification accuracy. Deployed at Amkor, South Korea, and trained the client team.

Senior Software Engineer - ML

Infosys | May 2016 - Jul 2019

Built an ML-powered lead-conversion prediction system and business-intelligence applications that moved sales prioritisation from intuition to data.

Achievements

15 honors and awards

Best Buy Talent Quotient (2026); Best Buy client AML Hall of Fame (2024), during the Wipro engagement.

Vehicle Routing & Transport Planning

Accenture | Google OR-Tools | Operations research

Designed vehicle-routing and mixed-integer models covering capacity, time windows, waiting time, demand, and departure constraints, then automated daily operational planning.

Industry Forecasting Platform

GD Research Centre | Time series | Scenario analysis

Designed and deployed domain-aware forecasting across 15+ industries, combining feature engineering, time-series modelling, and scenario analysis for enterprise decisions.

System Ownership

Design | Implementation | Production operations | Maintenance

Owned AI/ML systems from data and training pipelines through evaluation, deployment, monitoring, maintenance, incident support, and continuous improvement.

Engineering Leadership & Team Building

Architecture | Mentoring | Hiring | Cross-functional delivery

Stayed hands-on while leading ML/MLOps delivery, mentoring engineers, coordinating client and platform stakeholders, and helping scale the India AI engineering hub during the Wipro engagement through 300+ résumé screens and 50+ technical interviews.