

ENGINEERING

Python SQL Distributed Systems
API Services Async Processing

AI INFRASTRUCTURE

vLLM LoRA / QLoRA PyTorch RAPIDS
Morpheus Distributed Training

CLOUD & PLATFORMS

GCP Azure GKE Kubernetes Docker
Helm Istio Vertex AI Seldon Core

DEPLOYMENT & RELIABILITY

GitOps CI/CD Cloud Build ArgoCD
KCC Prometheus Grafana
Blue-Green Delivery JWT Security

CERTIFICATIONS

Oracle Cloud Infrastructure 2024
Generative AI Certified Professional
Microsoft Certified
Azure AI Engineer Associate
Microsoft Certified
Azure Data Scientist Associate
Google Cloud
Certified Generative AI Leader

EDUCATION

B.E. Electrical & Electronics
Andhra University | 2012 - 2016

ML METHODS

Scikit-learn TensorFlow Keras
DenseNet Classification
Predictive Modelling
Unsupervised Learning
Feature Engineering Transfer Learning
Anomaly Detection NLP
Time-Series Forecasting

AI infrastructure and MLOps engineer with 10+ years deploying enterprise ML across public cloud, production Kubernetes, and private GPU systems. At Best Buy, built shared LoRA/QLoRA training and vLLM serving on GKE, replacing external APIs with estimated annual API-cost savings of **~\$300K**. At Wipro, operated cybersecurity ML on a private NVIDIA DGX with 8xA100 GPUs and billion-row workloads. Brings hands-on Kubernetes, Helm, CI/CD, cloud, secure deployment, observability, troubleshooting, client delivery, and production ownership.

EXPERIENCE

Lead Artificial Intelligence Engineer
Best Buy India

Jul 2025 - Present | Bengaluru

- > Built and operated a shared production LLM platform covering config-driven LoRA/QLoRA training, automated evaluation, artifact lifecycle, model caching, and vLLM OpenAI-compatible serving on GKE.
- > Designed HPA autoscaling, Istio routing, blue-green releases, Cloud Build, ArgoCD, KCC, Prometheus, and Grafana; improved reliability, data governance, and cost with estimated annual API-cost savings of **~\$300K**. Kept model inputs company-managed and reduced external API exposure during peak demand.
- > Built fraud systems across batch, streaming, and real-time models. Beam/Dataflow feature pipelines handled **40,000 events/sec** at peak; supported a **200 ms scoring SLA** with secure, monitored GKE/Seldon Core inference and load testing.
- > Partnered with data science, platform, fraud, security, and leadership teams to assess requirements, review designs, troubleshoot delivery issues, and guide production handover.

Lead ML Engineer — Machine Learning & MLOps
Wipro

Aug 2023 - Jul 2025 | Hyderabad

- > Built distributed training and inference pipelines on a private NVIDIA DGX platform with **8xA100 GPUs**, processing billion-row AWS CloudTrail datasets using Morpheus, RAPIDS, continuous retraining, and Grafana alerting.
- > Engineered phishing classification at **2,400 emails/sec** and multiformat security automation across Outlook APIs, OCR, analysts, and external response partners.
- > Led and mentored ML/MLOps engineers; coordinated client/NVIDIA delivery; completed **300+ résumé screens** and conducted **50+ technical interviews**.

Senior Data Scientist — Data Science / Azure ML
Accenture

May 2021 - Sep 2023 | India

- > Led end-to-end Azure ML delivery for BamaGruppen, increasing inspections from **130 to 160 trucks/day**, improving low-quality detection 17%, and generating **\$2.5M annual savings**.
- > Translated client logistics needs into supported Python services with Azure Functions, Logic Apps, Azure DevOps, CI/CD, Google OR-Tools, stakeholder demos, alerts, troubleshooting, and measurable outcome tracking.

PROJECTS

Private LLM Serving Platform
Best Buy | GKE | vLLM | Safe rollout

Shared Python platform for training, evaluation, artifact management, autoscaled inference, traffic routing, blue-green release, security, and full-stack observability.

On-Prem GPU Security ML
Wipro | DGX | 8xA100 | Billion-row data

Distributed training and inference on private GPU infrastructure with high-throughput data processing, continuous retraining, monitoring, alerting, and analyst feedback.

GPU STACK

- CUDA GPU Libraries
- RAPIDS cuDF / CuPy
- Multi-GPU Training

OPTIMISATION & SIMULATION

- Google OR-Tools
- Vehicle Routing
- Mixed-Integer Programming
- Constraint Modelling
- Operations Research
- Assignment Models
- Supply-Chain Optimization
- Scenario Analysis

CLOUD & MANAGED ML

- GCS
- Artifact Registry
- Azure ML
- Azure DevOps
- Function Apps
- Logic Apps
- Google KCC

DATA & ML LIFECYCLE

- Apache Beam
- Dataflow
- Pub/Sub
- Vertex AI Feature Store
- Streaming Pipelines
- Training Pipelines
- Automated Evaluation
- Experiment Tracking
- Model Registry
- Model Versioning
- Continuous Training

LANGUAGES & DATA

- Pandas
- NumPy
- MS SQL Server
- IBM DB2
- Power BI

ENGINEERING LEADERSHIP

- Architecture Reviews
- Technical Mentoring
- Technical Hiring
- Team Building
- Stakeholder Alignment
- Cross-Functional Delivery

ACHIEVEMENTS

15 honors and awards across Best Buy, Wipro, Accenture, and Infosys.
Best Buy Talent Quotient (2026): cybersecurity and fraud delivery.
Best Buy client AML Hall of Fame (2024): DGX delivery during the Wipro engagement.

ADDITIONAL PROJECTS & EXPERIENCE

Real-Time Fraud Platform

Best Buy | 40k feature events/sec | 200 ms scoring SLA
Streaming features and low-latency inference on GCP and Kubernetes, with secure endpoints, load tests, telemetry, operational alerts, and GitOps delivery.

Client ML Delivery

Accenture | Azure | Python | CI/CD
Moved optimisation and quality-prediction systems from discovery to production services, automated pipelines, client demos, support, and measurable operating gains.

EARLIER EXPERIENCE & ACHIEVEMENTS

Data Scientist

GD Research Centre | May 2020 - May 2021
Built domain-aware time-series forecasting and scenario analysis across energy, manufacturing, healthcare, finance, retail, telecommunications, automotive, and other industry intelligence centers.

Deep Learning Engineer

Tata Consultancy Services | Jul 2019 - Mar 2020
Built and deployed an end-to-end semiconductor ML system: distributed training on ~1 million images in ~3 hours, with 90% classification accuracy. Deployed at Amkor, South Korea, and trained the client team.

Senior Software Engineer - ML

Infosys | May 2016 - Jul 2019
Built an ML-powered lead-conversion prediction system and business-intelligence applications that moved sales prioritisation from intuition to data.

Achievements

15 honors and awards
Best Buy Talent Quotient (2026); Best Buy client AML Hall of Fame (2024), during the Wipro engagement.

Vehicle Routing & Transport Planning

Accenture | Google OR-Tools | Operations research
Designed vehicle-routing and mixed-integer models covering capacity, time windows, waiting time, demand, and departure constraints, then automated daily operational planning.

Industry Forecasting Platform

GD Research Centre | Time series | Scenario analysis
Designed and deployed domain-aware forecasting across 15+ industries, combining feature engineering, time-series modelling, and scenario analysis for enterprise decisions.

System Ownership

Design | Implementation | Production operations | Maintenance
Owned AI/ML systems from data and training pipelines through evaluation, deployment, monitoring, maintenance, incident support, and continuous improvement.

Engineering Leadership & Team Building

Architecture | Mentoring | Hiring | Cross-functional delivery
Stayed hands-on while leading ML/MLOps delivery, mentoring engineers, coordinating client and platform stakeholders, and helping scale the India AI engineering hub during the Wipro engagement through 300+ résumé screens and 50+ technical interviews.