

PLATFORM ENGINEERING

Distributed Systems

Kubernetes

GKE

Docker

Helm

Istio

Python

AUTOMATION & DELIVERY

GitOps

ArgoCD

Cloud Build

Azure DevOps

Google KCC

CI/CD

Blue-Green Delivery

RELIABILITY & SECURITY

Prometheus

Grafana

Alerting

Load Testing

Incident Debugging

JWT Security

Production Handover

AI & GPU RUNTIME

vLLM

RAPIDS

MLflow

Seldon Core

Batch / Online Inference

CERTIFICATIONS

Oracle Cloud Infrastructure 2024
Generative AI Certified Professional

Microsoft Certified
Azure AI Engineer Associate

Microsoft Certified
Azure Data Scientist Associate

Google Cloud
Certified Generative AI Leader

EDUCATION

B.E. Electrical & Electronics
Andhra University | 2012 - 2016

ML METHODS

Scikit-learn

PyTorch

TensorFlow

Keras

DenseNet

Classification

Predictive Modelling

Unsupervised Learning

Feature Engineering

Transfer Learning

Anomaly Detection

NLP

Time-Series Forecasting

Hands-on AI platform and infrastructure engineer with 10+ years building high-availability systems across Kubernetes, distributed data, deployment automation, security, monitoring, and production support. At Best Buy, architected shared vLLM serving and LoRA/QLoRA training on GKE, with estimated annual API-cost savings of **~\$300K**, and leads real-time ML services with feature pipelines handling **40,000 events/sec** at peak and a **200 ms scoring SLA**. Earlier operated billion-row GPU security workloads on an 8xA100 DGX, delivered Azure ML systems saving \$2.5M annually, and deployed a multi-GPU vision system at a client site in South Korea. Strong in platform architecture, Kubernetes, GitOps, CI/CD, observability, capacity and performance thinking, incident debugging, technical standards, mentoring, and cross-team delivery.

EXPERIENCE

Lead Artificial Intelligence Engineer

Best Buy India

Jul 2025 - Present | Bengaluru

- > Architected a reusable production LLM platform for config-driven LoRA/QLoRA training, automated evaluation, model artifacts, and shared vLLM OpenAI-compatible serving on GKE.
- > Defined Kubernetes autoscaling, Istio routing, blue-green releases, Cloud Build, ArgoCD, Helm, Prometheus, and Grafana patterns; self-hosted inference supports estimated annual API-cost savings of **~\$300K** while strengthening operational control. Kept model inputs company-managed and reduced external API exposure during peak demand.
- > Built fraud systems across batch, streaming, and real-time models. Beam/Dataflow feature pipelines handled **40,000 events/sec** at peak; supported a **200 ms scoring SLA** with secure, monitored GKE/Seldon Core inference and load testing.
- > Lead technical direction across data science, platform, fraud, and security teams; review designs and mentor engineers.

Lead ML Engineer — Machine Learning & MLOps

Wipro

Aug 2023 - Jul 2025 | Hyderabad

- > Designed and operated GPU-accelerated AI systems on NVIDIA DGX (8xA100, 1TB RAM), running continuously refreshed training and inference across **billion-row** AWS CloudTrail datasets with Morpheus, RAPIDS, MLflow, and Grafana.
- > Engineered production phishing classification at **2,400 emails/sec** - 2x prior throughput - plus GPU-accelerated OCR and multiformat security pipelines integrated with Microsoft Sentinel.
- > Led and mentored ML/MLOps engineers; coordinated client/NVIDIA delivery; completed **300+ résumé screens** and conducted **50+ technical interviews**.

Senior Data Scientist — Data Science / Azure ML

Accenture

May 2021 - Sep 2023 | India

- > Led Azure ML platform delivery from requirements and architecture through automated retraining, deployment, monitoring, client demonstrations, and support; increased inspections from **130 to 160 trucks/day** and generated **\$2.5M annual savings**.
- > Built production Python services and delivery automation with Azure Functions, Logic Apps, Azure DevOps, Blob Storage, and CI/CD; translated complex operational needs into maintainable cloud systems.

PROJECTS

Enterprise AI Platform

Best Buy | GKE | vLLM | GitOps

Shared runtime services, autoscaling, traffic management, controlled rollout, security, cost optimization, observability, and reusable engineering standards.

Secure Real-Time Platform

Best Buy | 40k feature events/sec | 200 ms scoring SLA

Distributed pipelines, Kubernetes services, JWT security, load testing, telemetry, alerting, and reliability-focused delivery.

GPU STACK

- CUDA GPU Libraries
- RAPIDS cuDF / CuPy
- Multi-GPU Training

OPTIMISATION & SIMULATION

- Google OR-Tools
- Vehicle Routing
- Mixed-Integer Programming
- Constraint Modelling
- Operations Research
- Assignment Models
- Supply-Chain Optimization
- Scenario Analysis

CLOUD & MANAGED ML

- GCS
- Artifact Registry
- Azure ML
- Function Apps
- Logic Apps

DATA & ML LIFECYCLE

- Apache Beam
- Dataflow
- Pub/Sub
- Vertex AI Feature Store
- Streaming Pipelines
- Training Pipelines
- Automated Evaluation
- Experiment Tracking
- Model Registry
- Model Versioning
- Continuous Training

LANGUAGES & DATA

- SQL
- Pandas
- NumPy
- MS SQL Server
- IBM DB2
- Power BI

ENGINEERING LEADERSHIP

- Architecture Reviews
- Technical Mentoring
- Technical Hiring
- Team Building
- Stakeholder Alignment
- Cross-Functional Delivery

ACHIEVEMENTS

15 honors and awards across Best Buy, Wipro, Accenture, and Infosys.
Best Buy Talent Quotient (2026): cybersecurity and fraud delivery.
Best Buy client AML Hall of Fame (2024): DGX delivery during the Wipro engagement.

ADDITIONAL PROJECTS & EXPERIENCE

GPU Security Platform

Wipro | DGX 8xA100 | 1B+ rows

Training and inference operations, performance engineering, continuous refresh, monitoring, troubleshooting, and production handover.

On-Site Vision Deployment

TCS | Amkor | South Korea | 1M images

Built and deployed an end-to-end semiconductor ML system: distributed training on ~1 million images in ~3 hours, with 90% classification accuracy. Deployed at Amkor, South Korea, and trained the client team.

EARLIER EXPERIENCE & ACHIEVEMENTS

Data Scientist

GD Research Centre | May 2020 - May 2021

Built domain-aware time-series forecasting and scenario analysis across energy, manufacturing, healthcare, finance, retail, telecommunications, automotive, and other industry intelligence centers.

Deep Learning Engineer

Tata Consultancy Services | Jul 2019 - Mar 2020

Built and deployed an end-to-end semiconductor ML system: distributed training on ~1 million images in ~3 hours, with 90% classification accuracy. Deployed at Amkor, South Korea, and trained the client team.

Senior Software Engineer - ML

Infosys | May 2016 - Jul 2019

Built an ML-powered lead-conversion prediction system and business-intelligence applications that moved sales prioritisation from intuition to data.

Achievements

15 honors and awards

Best Buy Talent Quotient (2026); Best Buy client AML Hall of Fame (2024), during the Wipro engagement.

Vehicle Routing & Transport Planning

Accenture | Google OR-Tools | Operations research

Designed vehicle-routing and mixed-integer models covering capacity, time windows, waiting time, demand, and departure constraints, then automated daily operational planning.

Industry Forecasting Platform

GD Research Centre | Time series | Scenario analysis

Designed and deployed domain-aware forecasting across 15+ industries, combining feature engineering, time-series modelling, and scenario analysis for enterprise decisions.

System Ownership

Design | Implementation | Production operations | Maintenance

Owned AI/ML systems from data and training pipelines through evaluation, deployment, monitoring, maintenance, incident support, and continuous improvement.

Engineering Leadership & Team Building

Architecture | Mentoring | Hiring | Cross-functional delivery

Stayed hands-on while leading ML/MLOps delivery, mentoring engineers, coordinating client and platform stakeholders, and helping scale the India AI engineering hub during the Wipro engagement through 300+ résumé screens and 50+ technical interviews.