

PROGRAMMING & AI

PythonPyTorchTransformersPEFT

LoRA / QLoRAvLLMXGBoostSQL

AI PLATFORM

Model ServingOpenAI-Compatible APIs

Artifact LifecycleAutomated Evaluation

Batch / Online InferenceMLflow

Vertex AI PipelinesModel Registry

CLOUD & RUNTIME

GCPAzure MLGKEKubernetes

DockerDataflowSeldon Core

SECURITY & RELIABILITY

CI/CDArgoCDIstioHelm

Blue-Green DeliveryPrometheusGrafana

Load TestingJWT Security

CERTIFICATIONS

Oracle Cloud Infrastructure 2024
Generative AI Certified Professional

Microsoft Certified
Azure AI Engineer Associate

Microsoft Certified
Azure Data Scientist Associate

Google Cloud
Certified Generative AI Leader

EDUCATION

B.E. Electrical & Electronics
Andhra University | 2012 - 2016

ML METHODS

Scikit-learnTensorFlowKeras

DenseNetClassification

Predictive Modelling

Unsupervised Learning

Feature EngineeringTransfer Learning

Anomaly DetectionNLP

Time-Series Forecasting

AI platform and solutions engineer with 10+ years building, architecting, and operating production software, ML, and LLM systems across platform direction, serving, security, deployment, evaluation, and observability. At Best Buy, designed reusable LoRA/QLoRA training and shared vLLM serving on GKE, with estimated annual API-cost savings of **~\$300K**, and leads secure real-time fraud ML with feature pipelines handling **40,000 events/sec** at peak and a **200 ms scoring SLA**. Earlier built billion-row GPU pipelines on an 8xA100 DGX and Azure ML products saving \$2.5M annually. Hands-on with Python, Kubernetes, GCP, Azure ML, model APIs, CI/CD, controlled releases, load testing, monitoring, architecture reviews, technical standards, mentoring, and cross-team platform leadership.

EXPERIENCE

Lead Artificial Intelligence Engineer

Best Buy India
Jul 2025 - Present | Bengaluru

> Architected a reusable production LLM platform for config-driven LoRA/QLoRA training, automated post-training evaluation, artifact lifecycle, and vLLM OpenAI-compatible serving on GKE.

> Designed shared model APIs, Kubernetes autoscaling, Istio routing, blue-green releases, Cloud Build, ArgoCD, Helm, Prometheus, and Grafana; self-hosted inference supports estimated annual API-cost savings of **~\$300K** while strengthening data governance. Kept model inputs company-managed and reduced external API exposure during peak demand.

> Built fraud systems across batch, streaming, and real-time models. Beam/Dataflow feature pipelines handled **40,000 events/sec** at peak; supported a **200 ms scoring SLA** with secure, monitored GKE/Seldon Core inference and load testing.

> Lead technical direction across data science, platform, fraud, and security teams; review designs and mentor engineers.

Lead ML Engineer — Machine Learning & MLOps

Wipro
Aug 2023 - Jul 2025 | Hyderabad

> Built continuously retrained security-analytics pipelines on NVIDIA DGX (8xA100), processing **billion-row** AWS CloudTrail datasets with Morpheus, RAPIDS, MLflow, and Grafana.

> Engineered production phishing classifiers at **2,400 emails/sec** - 2x prior throughput - plus GPU-accelerated OCR and unstructured email pipelines integrated with Microsoft Sentinel.

> Led and mentored ML/MLOps engineers; coordinated client/NVIDIA delivery; completed **300+ résumé screens** and conducted **50+ technical interviews**.

Senior Data Scientist — Data Science / Azure ML

Accenture
May 2021 - Sep 2023 | India

> Led end-to-end Azure ML quality-prediction delivery, including retraining, deployment, monitoring and support; increased inspections from **130 to 160 trucks/day** and generated **\$2.5M annual savings**.

> Built tested Python optimisation and ML services with Azure Functions, Logic Apps, Azure DevOps and CI/CD; translated ambiguous client needs into reliable operated products.

PROJECTS

Shared LLM Platform Architecture

Best Buy | PyTorch | GKE | vLLM

Config-driven training, automated evaluation, model artifacts, shared OpenAI-compatible serving, autoscaling, controlled rollout, governance, and observability.

Secure Real-Time AI Architecture

Best Buy | 40k feature events/sec | 200 ms scoring SLA

Distributed feature engineering, registered models, Kubernetes inference, JWT-protected APIs, load testing, telemetry, and GitOps delivery.

GPU STACK

- CUDA GPU Libraries
- RAPIDS cuDF / CuPy
- Multi-GPU Training

OPTIMISATION & SIMULATION

- Google OR-Tools
- Vehicle Routing
- Mixed-Integer Programming
- Constraint Modelling
- Operations Research
- Assignment Models
- Supply-Chain Optimization
- Scenario Analysis

CLOUD & MANAGED ML

- GCS
- Artifact Registry
- Azure DevOps
- Function Apps
- Logic Apps
- Cloud Build
- Google KCC

DATA & ML LIFECYCLE

- Apache Beam
- Pub/Sub
- Vertex AI Feature Store
- Streaming Pipelines
- Training Pipelines
- Experiment Tracking
- Model Versioning
- Continuous Training

LANGUAGES & DATA

- Pandas
- NumPy
- MS SQL Server
- IBM DB2
- Power BI

ENGINEERING LEADERSHIP

- Architecture Reviews
- Technical Mentoring
- Technical Hiring
- Team Building
- Stakeholder Alignment
- Cross-Functional Delivery

ACHIEVEMENTS

15 honors and awards across Best Buy, Wipro, Accenture, and Infosys.
Best Buy Talent Quotient (2026): cybersecurity and fraud delivery.
Best Buy client AML Hall of Fame (2024): DGX delivery during the Wipro engagement.

ADDITIONAL PROJECTS & EXPERIENCE

GPU Security Analytics

Wipro | MLflow | DGX | Billion-row scale

Per-entity autoencoders and production inference across high-cardinality security datasets with GPU acceleration, monitoring, and continuous retraining.

Azure ML Quality Platform

Accenture | Azure ML | \$2.5M/year

Model optimisation, retraining automation, deployment, monitoring, client demonstrations, support, and measurable customer outcomes.

EARLIER EXPERIENCE & ACHIEVEMENTS

Data Scientist

GD Research Centre | May 2020 - May 2021

Built domain-aware time-series forecasting and scenario analysis across energy, manufacturing, healthcare, finance, retail, telecommunications, automotive, and other industry intelligence centers.

Deep Learning Engineer

Tata Consultancy Services | Jul 2019 - Mar 2020

Built and deployed an end-to-end semiconductor ML system: distributed training on ~1 million images in ~3 hours, with 90% classification accuracy. Deployed at Amkor, South Korea, and trained the client team.

Senior Software Engineer - ML

Infosys | May 2016 - Jul 2019

Built an ML-powered lead-conversion prediction system and business-intelligence applications that moved sales prioritisation from intuition to data.

Achievements

15 honors and awards

Best Buy Talent Quotient (2026); Best Buy client AML Hall of Fame (2024), during the Wipro engagement.

Vehicle Routing & Transport Planning

Accenture | Google OR-Tools | Operations research

Designed vehicle-routing and mixed-integer models covering capacity, time windows, waiting time, demand, and departure constraints, then automated daily operational planning.

Industry Forecasting Platform

GD Research Centre | Time series | Scenario analysis

Designed and deployed domain-aware forecasting across 15+ industries, combining feature engineering, time-series modelling, and scenario analysis for enterprise decisions.

System Ownership

Design | Implementation | Production operations | Maintenance

Owned AI/ML systems from data and training pipelines through evaluation, deployment, monitoring, maintenance, incident support, and continuous improvement.

Engineering Leadership & Team Building

Architecture | Mentoring | Hiring | Cross-functional delivery

Stayed hands-on while leading ML/MLOps delivery, mentoring engineers, coordinating client and platform stakeholders, and helping scale the India AI engineering hub during the Wipro engagement through 300+ résumé screens and 50+ technical interviews.