

LLM & GENAI SYSTEMS

LLMsLoRA / QLoRAPEFT

UnslothvLLMTransformers

Fine-tuningModel Evaluation

LLM Deployment

ML & DEEP LEARNING

PyTorchTensorFlow

Scikit-learnNumPyPandas

XGBoostRAPIDS (cuDF/CuPy)

Computer VisionAutoencoders

Distributed Training

MLOPS & PLATFORMS

Cloud & Orchestration

GCP Vertex AIGKEDataflow

GCSAzure Machine Learning

Azure Function Apps

Deployment & MLOps

KubernetesHelmIstio

DockerSeldon CoreMLflow

Model Registry

Observability & CI/CD

GrafanaPrometheus

Cloud BuildArgoCD

Azure DevOps

LANGUAGES & ARCHITECTURE

PythonSQL

Distributed Systems

Software Architecture

Real-time Streaming

Testing & Debugging

CERTIFICATIONS

Oracle Cloud Infrastructure 2024
Generative AI Certified Professional

Microsoft Certified
Azure AI Engineer Associate

Microsoft Certified
Azure Data Scientist Associate

Lead ML engineer with 10+ years of engineering experience. Build reusable LLM fine-tuning, evaluation, and vLLM serving with estimated annual API-cost savings of **~\$300K**, high-throughput fraud feature pipelines, and billion-row GPU cybersecurity systems. Lead architecture, model lifecycle, cloud delivery, monitoring, and mentoring; bring hands-on Python and client-facing production experience.

WORK EXPERIENCE

- Lead Artificial Intelligence Engineer

Best Buy · Bengaluru, India

Jul 2025 - Present

 - Architected a config-driven Python LoRA/QLoRA fine-tuning and evaluation framework using Unsloth, PEFT, Transformers, and Accelerate; containerized vLLM on GKE with autoscaling, Istio, blue-green releases, and observability supported estimated annual API-cost savings of **~\$300K**. Kept model inputs company-managed and reduced external API exposure during peak demand.
 - Built fraud systems across batch, streaming, and real-time models. Beam/Dataflow feature pipelines handled **40,000 events/sec** at peak; supported a **200 ms scoring SLA** with secure, monitored GKE/Seldon Core inference and load testing.
 - Lead architecture and production delivery across search, personalization, fraud, and cybersecurity; partner with data scientists to turn validated models into secure, monitored services.
- Lead ML Engineer - ML Engineering & Operations

Wipro · Client: Major US Retail Corporation

Aug 2023 - Jul 2025

 - Led ML/MLOps engineers and Morpheus/RAPIDS model lifecycles on **NVIDIA DGX (8×A100)**, training per-entity autoencoders on **billion-row CloudTrail datasets** with MLflow model management, continuous evaluation, debugging, and Grafana alerting. Completed **300+ résumé screens** and **50+ technical interviews** at Wipro.
 - Launched phishing classification at **2,400 emails/sec - 2× prior throughput** and OCR-based abuse-inbox triage that doubled case volume, integrating predictions with Microsoft Sentinel and security-response workflows.
- Senior Data Scientist — Data Science / Azure ML

Accenture

May 2021 - Sep 2023

 - Led cross-functional architecture and launch of an Azure ML quality-prediction product: **\$2.5M annual savings**, inspections **130 → 160/day**, **+17%** defect detection, and **+20%** customer approval.
 - Operationalized data preparation, optimization, retraining, monitoring, and CI/CD with Azure ML Studio, Function Apps, Logic Apps, and Azure DevOps; mentored engineers and coordinated business, data science, platform, and client stakeholders.

EARLIER EXPERIENCE

Data Scientist

GD Research Centre Pvt Ltd

May 2020 - May 2021

Designed and deployed Python forecasting pipelines across 15+ industry verticals, from heterogeneous data preparation through client-facing results.

Deep Learning Engineer

Tata Consultancy Services

Jul 2019 - Mar 2020

Built and deployed an end-to-end semiconductor ML system: distributed training on **~1 million images in ~3 hours; 90% classification accuracy**. Deployed at Amkor, South Korea, and trained the client team.

Senior Software Engineer - ML

Infosys

May 2016 - Jul 2019

Delivered design, coding, unit/integration testing, and production support; transitioned to Python ML in 2018 and shipped lead-conversion prediction.

EDUCATION

B.E. Electrical & Electronics

Andhra University, Visakhapatnam
2012 - 2016 · 80%

GPU & DISTRIBUTED
COMPUTING

CUDA GPU Libraries

GPU-Accelerated Computing

OPTIMISATION & SIMULATION

Google OR-Tools

Vehicle Routing

Mixed-Integer Programming

Constraint Modelling

Operations Research

Supply-Chain Optimization

MODEL LIFECYCLE

Experiment Tracking

Automated Retraining

Vertex AI Pipelines

Azure ML Pipelines

Vertex AI Feature Store

CLOUD SERVICES

Cloud Run

Cloud Functions

Logic Apps

Google KCC

PRODUCTION ENGINEERING

ML System Architecture

Cloud-Native ML Systems

GitOps

JWT Authentication

Load Testing

Unit / Integration Testing

LANGUAGES & DATA

BigQuery

MS SQL Server

Redis Caching

Azure Blob Storage

Artifact Registry

Dask

Git

Linux / Shell Scripting

ENGINEERING LEADERSHIP

Architecture Reviews

Technical Mentoring

Technical Hiring

Team Building

Stakeholder Alignment

Cross-Functional Delivery

PROJECTS

LLM Fine-tuning & Serving Framework

Best Buy · Jul 2025 - Present

Config-driven Python platform for LoRA/QLoRA fine-tuning, data preparation, multi-task training, automated evaluation, artifact management, and vLLM serving. Supports Unsloth and PEFT/Transformers with Accelerate; adapters hot-load into a shared GKE deployment.

~\$300K estimated annual API-cost savings

AWS Account Behavioral Fingerprinting

Wipro · Jan 2024 - Aug 2025

Per-entity autoencoders learn normal AWS account, user, service, and machine activity. Morpheus/RAPIDS training and inference run on DGX (8×A100) over billion-row CloudTrail datasets with continuous retraining and Grafana alerts.

RAPIDS / cuDF

Billion-row datasets · Continuous model lifecycle · 8×A100 DGX

Proofpoint Phishing Classification

Wipro · Oct 2023 - Feb 2024

Built Morpheus training and inference pipelines on DGX for preprocessing, feature engineering, model evaluation, and live scoring. Integrated security-team feedback and Microsoft Sentinel data to catch phishing that bypassed initial filters.

RAPIDS

DGX

Microsoft Sentinel

2,400 emails/sec · 2× prior throughput · Production security workflow

Real-time Retail Fraud ML Platform

Best Buy · Aug 2025 - Jan 2026

Large-scale Beam/Dataflow pipelines transform clickstream and transaction events into Vertex AI Feature Store signals. XGBoost models deploy through Seldon Core on GKE with JWT security, load testing, GitOps releases, and production monitoring.

Apache Beam

Vertex AI

40k feature events/sec · 200 ms scoring SLA · Full lifecycle deployment

Semiconductor Computer Vision Platform

TCS / Amkor Technology · Jul 2019 - Mar 2020

Built and deployed an end-to-end semiconductor ML system: distributed training on ~1 million images in ~3 hours, with 90% classification accuracy. Deployed at Amkor, South Korea, and trained the client team.

DenseNet

Distributed GPU

~1M images · ~3-hour training cycles · Production launch in South Korea

Azure ML Product Quality Platform

Accenture / BamaGruppen · May 2021 - Dec 2021

Led an end-to-end ML product from data preparation and model optimization through automated retraining, CI/CD, monitoring, client demonstrations, and ongoing production support on Azure.

Azure ML

Logic Apps

CI/CD

Monitoring

\$2.5M/year · +17% detection · +20% approval · Launched product

LEADERSHIP & CONSULTING DELIVERY

Technical Direction

Architecture & Engineering Standards

Sets priorities and design direction for production ML platforms spanning fine-tuning, evaluation, serving, streaming data, and observability.

People & Technical Leadership

Distributed, Matrixed Teams

Lead and mentors engineers, aligns data science, platform, security, fraud, business, and client teams, and translates ambiguous goals into delivery plans.

Consulting Delivery

Design Through Production

Translates stakeholder requirements into architectures, prototypes, production launches, monitoring, and support across retail, cybersecurity, supply chain, and semiconductor domains.