

Laxmi Narasimha Hari Yelesetty

Lead AI/ML Engineer & Data Scientist | Production ML & GenAI | Distributed Systems | GCP · Azure · GPU Computing

Bengaluru, India
+91 99161 12881 · ylnharimailme@gmail.com
linkedin.com/in/ylnhari · github.com/ylnhari
ylnhari.github.io

CORE EXPERTISE

ML systems

- ML infrastructure
- Model deployment
- Model evaluation
- Fine-tuning
- Classification
- Feature engineering
- Distributed training
- Testing & debugging

GenAI & serving

- LLMs
- LoRA / QLoRA
- PEFT
- Unsloth
- vLLM
- Transformers

Cloud & platforms

- GCP
- Vertex AI
- GKE
- Dataflow
- Pub/Sub
- Azure ML

Delivery & reliability

- Kubernetes
- Docker
- Helm
- Istio
- ArgoCD
- Cloud Build
- Prometheus
- Grafana

Languages

- Python
- SQL

CAREER EVIDENCE

10+ years

Software engineering and production delivery

5+ years

ML design, infrastructure, testing, launch and technical leadership

40k feature events/sec

Feature throughput; expected 200 ms scoring SLA

~\$300K estimated annual

API-cost savings

Self-hosted open-source LLM platform

Billion-row training

GPU-accelerated behavioral ML on 8xA100 DGX

Lead AI and ML engineer and data scientist with **10+ years** of experience, building production AI/ML/LLM systems. Leads end-to-end AI delivery, combining hands-on model development with technical leadership to turn business problems into scalable, high-throughput, low-latency production systems. Experience spans predictive modelling, forecasting, operations research, computer vision, and generative AI across retail, cybersecurity, manufacturing, insurance, and supply-chain optimization.

PROFESSIONAL EXPERIENCE

Lead Artificial Intelligence Engineer

Best Buy · Bengaluru, India

Jul 2025 – Present

- Architected a config-driven Python platform for LoRA/QLoRA fine-tuning, data preparation and automated evaluation across classification, information extraction and generation tasks; served PEFT adapters through vLLM on GKE with autoscaling, Istio routing and blue-green releases, with estimated annual API-cost savings of **~\$300K**. Kept model inputs company-managed and reduced external API exposure during peak demand.
- Lead end-to-end MLOps for search, fraud and security models using GCP, Vertex AI, GKE, Dataflow and Seldon Core, translating validated models into scalable services with production SLAs.
- Built fraud systems across batch, streaming, and real-time models. Beam/Dataflow feature pipelines handled **40,000 events/sec** at peak; supported a **200 ms scoring SLA** with secure, monitored GKE/Seldon Core inference and load testing.
- Lead technical direction across data science, platform, fraud, and security teams; review designs and mentor engineers.

Lead ML Engineer — Machine Learning & MLOps

Wipro · US retail and cybersecurity portfolio

Aug 2023 – Jul 2025

- Led ML/MLOps engineers building Morpheus and RAPIDS training and inference pipelines on NVIDIA DGX (8xA100), including per-entity autoencoders trained on billion-row CloudTrail datasets with continuous retraining, evaluation and Grafana alerting. Completed **300+ résumé screens** and **50+ technical interviews** at Wipro.
- Launched phishing classification at **2,400 emails/sec**—twice the prior throughput—and integrated scores with Microsoft Sentinel and incident-response workflows.
- Automated abuse-inbox classification, routing and OCR extraction of threat identifiers, doubling supported case volume and reducing manual security review.

Senior Data Scientist — Data Science / Azure ML

Accenture · Norway logistics portfolio

May 2021 – Sep 2023

- Led architecture and launch of an Azure ML quality-prediction product from data preparation and model optimization through automated retraining, testing, monitoring and support; delivered **\$2.5M annual savings**, 130→160 inspections/day, +17% defect detection and +20% customer approval.
- Designed Python and OR-Tools optimization services with Azure Functions, Logic Apps and CI/CD; mentored engineers and coordinated client, business and engineering stakeholders.

EARLIER EXPERIENCE

Data Scientist GD Research Centre

May 2020 – May 2021

Designed and deployed Python time-series forecasting pipelines across 15+ industry verticals.

Deep Learning Engineer Tata Consultancy Services

Jul 2019 – Mar 2020

Built and deployed an end-to-end semiconductor ML system: distributed training on **~1 million images in ~3 hours**, with **90% classification accuracy**. Deployed at Amkor, South Korea, and trained the client team.

Senior Software Engineer — ML Infosys

May 2016 – Jul 2019

Delivered design, code, unit/integration testing and production support; shipped lead-conversion prediction.

PROJECTS

LLM fine-tuning, evaluation & serving platform

Best Buy

Config-driven LoRA/QLoRA framework using Unsloth, PEFT, Transformers and Accelerate. Supports multi-task data preparation, automated evaluation, artifact versioning and hot-loaded vLLM adapters on a shared GKE service.

Python · Transformers · PEFT · Unsloth · vLLM · GKE · GCS · Prometheus

~\$300K estimated annual API-cost savings

Real-time retail fraud ML platform

Best Buy

Apache Beam and Dataflow pipelines transform clickstream and transaction events into Vertex AI feature signals. XGBoost models deploy through Seldon Core on GKE with JWT security, load testing, observability and controlled releases.

Beam · Dataflow · Vertex AI · GKE · Seldon Core · GitOps

40k feature events/sec · 200 ms scoring SLA

AWS behavioral fingerprinting

Wipro

Per-entity autoencoders learn normal account, user, service and machine behavior. GPU-accelerated training and inference process high-cardinality CloudTrail data with continuous retraining and live security alerts.

NVIDIA Morpheus · RAPIDS · cuDF · Autoencoders · MLflow · Grafana

Billion-row datasets · 8xA100 DGX

Phishing and abuse classification

Wipro

Production classification pipelines combine message preprocessing, feature engineering, model evaluation, feedback and inference. OCR and rules extract threat identifiers from varied email bodies and attachments for downstream action.

Python · Morpheus · RAPIDS · OCR · Microsoft Sentinel

2,400 emails/sec · 2x prior throughput

Semiconductor defect classification

TCS / Amkor Technology

Built and deployed an end-to-end semiconductor ML system: distributed training on ~1 million images in ~3 hours, with 90% classification accuracy. Deployed at Amkor, South Korea, and trained the client team.

Python · TensorFlow · DenseNet · Distributed GPU · Docker

~1M images · ~3-hour training cycle · Production launch

Product quality prediction

Accenture / BamaGruppen

Led the full ML product lifecycle on Azure: data preparation, modeling, automated retraining, CI/CD, monitoring, client demonstrations and production support. Partnered with inspection and business teams to measure impact.

Python · Azure ML · Logic Apps · CI/CD · Monitoring

\$2.5M/year · +17% detection · +20% approval

SYSTEM OWNERSHIP & ENGINEERING LEADERSHIP

Architecture & technical direction

Defines end-to-end ML system designs spanning data, training, evaluation, serving, deployment and observability. Makes trade-offs for latency, throughput, reliability and cost.

Cross-functional delivery

Aligns data science, platform, security, fraud, operations and client stakeholders; converts ambiguous goals into testable systems and measurable releases.

Engineering leadership

Lead and mentors ML/MLOps engineers, reviews designs and delivery, handles production issues, and helps scale teams through structured technical hiring.

CERTIFICATIONS & EDUCATION

Generative AI

Oracle Cloud Infrastructure

Generative AI Certified Professional

Cloud AI & ML

Google Cloud · Microsoft

Google Cloud Gen AI L3 · Azure AI Engineer Associate · Azure Data Scientist Associate

B.E., Electrical & Electronics

Andhra University

2012–2016

ADDITIONAL RECOGNITION

- › Best Buy Leadership — Talent Quotient recognition for production impact, reliability improvements, rapid GCP ramp-up and the self-hosted LLM cost initiative.
- › Wipro Above & Beyond, Pinnacle and Being Responsive awards for AI/ML delivery, technical leadership and hiring contributions.
- › Infosys Insta Awards for automated impact analysis and continuous defect-free delivery.