

SYSTEMS & AUTOMATION

- PythonLinux ShellSQL
- Distributed SystemsAPI Services
- Stream / Batch

ML LIFECYCLE & SERVING

- Azure MLVertex AI Pipelines
- Model RegistryFeature StoreMLflow
- vLLMSeldon CoreDistributed Training

CLOUD & ORCHESTRATION

- Azure DevOpsAzure FunctionsGCP
- GKEKubernetesDockerHelmIstio
- ArgoCDGoogle KCC

DELIVERY & RELIABILITY

- CI/CDGitOpsBlue-Green Delivery
- HPA AutoscalingPrometheusGrafana
- Load TestingJWT Security
- SLA Operations

CERTIFICATIONS

- Oracle Cloud Infrastructure 2024
Generative AI Certified Professional
- Microsoft Certified
Azure AI Engineer Associate
- Microsoft Certified
Azure Data Scientist Associate
- Google Cloud
Certified Generative AI Leader

EDUCATION

B.E. Electrical & Electronics
Andhra University | 2012 - 2016

ML METHODS

- Scikit-learnPyTorchTensorFlow
- KerasDenseNetClassification
- Predictive Modelling
- Unsupervised Learning
- Feature EngineeringTransfer Learning
- Anomaly DetectionNLP
- Time-Series Forecasting

Production MLOps and AI-platform engineer with 10+ years building reliable software and end-to-end ML systems across Azure, GCP, Kubernetes, and on-prem GPU infrastructure. At Best Buy, architected shared training, automated evaluation, artifact lifecycle, and vLLM serving on GKE, with estimated annual API-cost savings of **~\$300K**; also operates real-time ML infrastructure with feature pipelines handling **40,000 events/sec** at peak and a **200 ms scoring SLA**. Earlier built continuously retrained billion-row pipelines on an 8xA100 DGX and led Azure ML delivery generating \$2.5M annual savings. Hands-on with Python, model serving, lifecycle automation, CI/CD, controlled releases, observability, security, and production debugging.

EXPERIENCE

Lead Artificial Intelligence Engineer

Best Buy India

Jul 2025 - Present | Bengaluru

- > Architected a shared production platform for config-driven LoRA/QLoRA training, automated post-training evaluation, artifact lifecycle, and vLLM OpenAI-compatible model serving on GKE.
- > Designed autoscaled container delivery with HPA, Istio traffic routing, blue-green releases, Cloud Build, ArgoCD, Google KCC, Prometheus, and Grafana; self-hosted services support estimated annual API-cost savings of **~\$300K**. Kept model inputs company-managed and reduced external API exposure during peak demand.
- > Built fraud systems across batch, streaming, and real-time models. Beam/Dataflow feature pipelines handled **40,000 events/sec** at peak; supported a **200 ms scoring SLA** with secure, monitored GKE/Seldon Core inference and load testing.
- > Lead technical direction across data science, platform, fraud, and security teams; review designs and mentor engineers.

Lead ML Engineer — Machine Learning & MLOps

Wipro

Aug 2023 - Jul 2025 | Hyderabad

- > Built distributed training and inference pipelines on NVIDIA DGX (8xA100), processing **billion-row** AWS CloudTrail datasets with Morpheus, RAPIDS, MLflow model management, continuous retraining, and Grafana alerting.
- > Engineered phishing classification at **2,400 emails/sec** - 2x prior throughput - and multiformat abuse-inbox processing integrated with Microsoft Sentinel and production security workflows.
- > Led and mentored ML/MLOps engineers; coordinated client/NVIDIA delivery; completed **300+ résumé screens** and conducted **50+ technical interviews**.

Senior Data Scientist — Data Science / Azure ML

Accenture

May 2021 - Sep 2023 | India

- > Led end-to-end Azure ML product delivery, including automated retraining, deployment, monitoring, and support; increased inspections from **130 to 160 trucks/day** and generated **\$2.5M annual savings**.
- > Built production services and CI/CD automation with Python, Azure Functions, Logic Apps, Azure DevOps, Azure storage, and Google OR-Tools; translated ambiguous client needs into operated systems.

PROJECTS

End-to-End ML Platform

Best Buy | Vertex AI | Seldon

Streaming features, model registry and pipelines, Kubernetes serving, GitOps delivery, secure endpoints, load testing, telemetry, and a 200ms SLA.

LLM Training & Serving

Best Buy | GKE | vLLM | Evaluation

Shared training, automated evaluation, artifact management, model caching, autoscaled inference, traffic routing, controlled rollout, and observability.

GPU STACK

- CUDA GPU Libraries
- RAPIDS cuDF / CuPy
- Multi-GPU Training

OPTIMISATION & SIMULATION

- Google OR-Tools
- Vehicle Routing
- Mixed-Integer Programming
- Constraint Modelling
- Operations Research
- Assignment Models
- Supply-Chain Optimization
- Scenario Analysis

CLOUD & MANAGED ML

- GCS
- Artifact Registry
- Function Apps
- Logic Apps
- Cloud Build

DATA & ML LIFECYCLE

- Apache Beam
- Dataflow
- Pub/Sub
- Vertex AI Feature Store
- Streaming Pipelines
- Training Pipelines
- Automated Evaluation
- Experiment Tracking
- Model Versioning
- Continuous Training

LANGUAGES & DATA

- Pandas
- NumPy
- MS SQL Server
- IBM DB2
- Power BI

ENGINEERING LEADERSHIP

- Architecture Reviews
- Technical Mentoring
- Technical Hiring
- Team Building
- Stakeholder Alignment
- Cross-Functional Delivery

ACHIEVEMENTS

15 honors and awards across Best Buy, Wipro, Accenture, and Infosys.
Best Buy Talent Quotient (2026): cybersecurity and fraud delivery.
Best Buy client AML Hall of Fame (2024): DGX delivery during the Wipro engagement.

ADDITIONAL PROJECTS & EXPERIENCE

Cybersecurity GPU Pipelines

Wipro | MLflow | DGX | Billion-row scale

Continuously retrained behavioural models and high-throughput classifiers with GPU acceleration, lifecycle tracking, monitoring, and security integrations.

Azure ML Product Delivery

Accenture | Azure ML | \$2.5M/year

Production ML delivery spanning problem framing, retraining, automated deployment, monitoring, client handover, support, and measurable operational outcomes.

EARLIER EXPERIENCE & ACHIEVEMENTS

Data Scientist

GD Research Centre | May 2020 - May 2021

Built domain-aware time-series forecasting and scenario analysis across energy, manufacturing, healthcare, finance, retail, telecommunications, automotive, and other industry intelligence centers.

Deep Learning Engineer

Tata Consultancy Services | Jul 2019 - Mar 2020

Built and deployed an end-to-end semiconductor ML system: distributed training on ~1 million images in ~3 hours, with 90% classification accuracy. Deployed at Amkor, South Korea, and trained the client team.

Senior Software Engineer - ML

Infosys | May 2016 - Jul 2019

Built an ML-powered lead-conversion prediction system and business-intelligence applications that moved sales prioritisation from intuition to data.

Achievements

15 honors and awards

Best Buy Talent Quotient (2026); Best Buy client AML Hall of Fame (2024), during the Wipro engagement.

Vehicle Routing & Transport Planning

Accenture | Google OR-Tools | Operations research

Designed vehicle-routing and mixed-integer models covering capacity, time windows, waiting time, demand, and departure constraints, then automated daily operational planning.

Industry Forecasting Platform

GD Research Centre | Time series | Scenario analysis

Designed and deployed domain-aware forecasting across 15+ industries, combining feature engineering, time-series modelling, and scenario analysis for enterprise decisions.

System Ownership

Design | Implementation | Production operations | Maintenance

Owned AI/ML systems from data and training pipelines through evaluation, deployment, monitoring, maintenance, incident support, and continuous improvement.

Engineering Leadership & Team Building

Architecture | Mentoring | Hiring | Cross-functional delivery

Stayed hands-on while leading ML/MLOps delivery, mentoring engineers, coordinating client and platform stakeholders, and helping scale the India AI engineering hub during the Wipro engagement through 300+ résumé screens and 50+ technical interviews.