

PLATFORM ARCHITECTURE

- GCPGKEVertex AIDataflow
- Pub/SubAzure MLAzure Functions
- KubernetesDockerHelmIstio

AI & ML PLATFORMS

- LLMs vLLM LoRA / QLoRA PEFT
- Unsloth Transformers Seldon Core
- MLflow

RELIABILITY & DELIVERY

- Cloud BuildArgoCDGoogle KCC
- GitOpsCI/CDPrometheusGrafana
- Load TestingJWT Security
- Blue-Green Delivery

ENGINEERING

- PythonSQLDistributed Systems
- Real-Time StreamingGPU Computing
- Technical Leadership

CERTIFICATIONS

- Oracle Cloud Infrastructure 2024
Generative AI Certified Professional
- Microsoft Certified
Azure AI Engineer Associate
- Microsoft Certified
Azure Data Scientist Associate
- Google Cloud
Certified Generative AI Leader

EDUCATION

- B.E. Electrical & Electronics
Andhra University | 2012 - 2016

ML METHODS

- Scikit-learnPyTorchTensorFlow
- KerasDenseNetClassification
- Predictive Modelling
- Unsupervised Learning
- Feature EngineeringTransfer Learning
- Anomaly DetectionNLP
- Time-Series Forecasting

Hands-on AI platform engineer with 10+ years building and operating production ML systems across GCP, Azure, and on-prem GPU infrastructure. At Best Buy, architected LLM fine-tuning, evaluation, and vLLM serving on GKE, replacing external APIs with estimated annual API-cost savings of **~\$300K**; also owns real-time ML infrastructure with feature pipelines handling **40,000 events/sec** at peak and a **200 ms scoring SLA**. Earlier led billion-row cybersecurity ML on 8xA100 DGX and an Azure ML product delivering \$2.5M annual savings. Combines cloud architecture, MLOps, CI/CD, GitOps, observability, security controls, cost ownership, and senior stakeholder delivery.

EXPERIENCE

Lead Artificial Intelligence Engineer

Best Buy India
Jul 2025 - Present | Bengaluru

- Architected a production LLM platform spanning config-driven LoRA/QLoRA training, automated evaluation, artifact management, and shared vLLM OpenAI-compatible serving on GKE.
- Designed the cloud delivery architecture with HPA autoscaling, Istio routing, blue-green releases, Cloud Build, ArgoCD, KCC, Prometheus, and Grafana; self-hosted models support estimated annual API-cost savings of **~\$300K**. Kept model inputs company-managed and reduced external API exposure during peak demand.
- Built fraud systems across batch, streaming, and real-time models. Beam/Dataflow feature pipelines handled **40,000 events/sec** at peak; supported a **200 ms scoring SLA** with secure, monitored GKE/Seldon Core inference and load testing.
- Lead technical direction across data science, platform, fraud, and security teams; review designs and mentor engineers.

Lead ML Engineer — Machine Learning & MLOps

Wipro
Aug 2023 - Jul 2025 | Hyderabad

- Led GPU-accelerated platform delivery on NVIDIA DGX (8xA100), training continuously refreshed per-entity autoencoders over **billion-row** AWS CloudTrail datasets with Morpheus, RAPIDS, MLflow, and Grafana.
- Built phishing training and inference pipelines at **2,400 emails/sec** - 2x prior throughput - and automated multiformat abuse-inbox triage for production security-response workflows.
- Led and mentored ML/MLOps engineers; coordinated client/NVIDIA delivery; completed **300+ résumé screens** and conducted **50+ technical interviews**.

Senior Data Scientist — Data Science / Azure ML

Accenture
May 2021 - Sep 2023 | India

- Led end-to-end Azure ML product delivery for BamaGruppen, increasing inspections from **130 to 160 trucks/day**, improving low-quality detection 17%, and generating **\$2.5M annual savings**.
- Built cloud-hosted optimization and simulation platforms using Python, Azure Functions, Logic Apps, Azure DevOps, CI/CD, and Google OR-Tools; translated client requirements into operated production services.

PROJECTS

LLM Training, Evaluation & Serving Platform

Best Buy | GKE | vLLM | GitOps

Config-driven Python platform for multi-task LoRA/QLoRA training, evaluation, artifact lifecycle, model caching, autoscaled serving, traffic management, and full-stack observability.

Real-Time Retail ML Platform

Best Buy | 40k feature events/sec | 200 ms scoring SLA

Streaming feature engineering and real-time inference across payment, refund, reship, chargeback, and item-not-received flows with cloud-native delivery and operational controls.

GPU STACK

- CUDA GPU Libraries
- RAPIDS cuDF / CuPy
- Multi-GPU Training

OPTIMISATION & SIMULATION

- Google OR-Tools
- Vehicle Routing
- Mixed-Integer Programming
- Constraint Modelling
- Operations Research
- Assignment Models
- Supply-Chain Optimization
- Scenario Analysis

CLOUD & MANAGED ML

- GCS
- Artifact Registry
- Azure DevOps
- Function Apps
- Logic Apps

DATA & ML LIFECYCLE

- Apache Beam
- Vertex AI Feature Store
- Streaming Pipelines
- Training Pipelines
- Automated Evaluation
- Experiment Tracking
- Model Registry
- Model Versioning
- Continuous Training

LANGUAGES & DATA

- Pandas
- NumPy
- MS SQL Server
- IBM DB2
- Power BI

ENGINEERING LEADERSHIP

- Architecture Reviews
- Technical Mentoring
- Technical Hiring
- Team Building
- Stakeholder Alignment
- Cross-Functional Delivery

ACHIEVEMENTS

15 honors and awards across Best Buy, Wipro, Accenture, and Infosys.
Best Buy Talent Quotient (2026): cybersecurity and fraud delivery.
Best Buy client AML Hall of Fame (2024): DGX delivery during the Wipro engagement.

ADDITIONAL PROJECTS & EXPERIENCE

Cybersecurity GPU Platform

Wipro | DGX | Billion-row scale

Morpheus and RAPIDS pipelines for behavioral anomaly models, continuous retraining, evaluation, alerting, and analyst-facing production workflows across cloud security logs.

Rover - Go Systems Project

Open source | Go | Real-time SSE

Built a zero-dependency remote command executor and project launcher in Go with persistent sessions and real-time server-sent event streaming.

EARLIER EXPERIENCE & ACHIEVEMENTS

Data Scientist

GD Research Centre | May 2020 - May 2021

Built domain-aware time-series forecasting and scenario analysis across energy, manufacturing, healthcare, finance, retail, telecommunications, automotive, and other industry intelligence centers.

Deep Learning Engineer

Tata Consultancy Services | Jul 2019 - Mar 2020

Built and deployed an end-to-end semiconductor ML system: distributed training on **~1 million images in ~3 hours**, with **90% classification accuracy**. Deployed at Amkor, South Korea, and trained the client team.

Senior Software Engineer - ML

Infosys | May 2016 - Jul 2019

Built an ML-powered lead-conversion prediction system and business-intelligence applications that moved sales prioritisation from intuition to data.

Achievements

15 honors and awards

Best Buy Talent Quotient (2026); Best Buy client AML Hall of Fame (2024), during the Wipro engagement.

Vehicle Routing & Transport Planning

Accenture | Google OR-Tools | Operations research

Designed vehicle-routing and mixed-integer models covering capacity, time windows, waiting time, demand, and departure constraints, then automated daily operational planning.

Industry Forecasting Platform

GD Research Centre | Time series | Scenario analysis

Designed and deployed domain-aware forecasting across 15+ industries, combining feature engineering, time-series modelling, and scenario analysis for enterprise decisions.

System Ownership

Design | Implementation | Production operations | Maintenance

Owned AI/ML systems from data and training pipelines through evaluation, deployment, monitoring, maintenance, incident support, and continuous improvement.

Engineering Leadership & Team Building

Architecture | Mentoring | Hiring | Cross-functional delivery

Stayed hands-on while leading ML/MLOps delivery, mentoring engineers, coordinating client and platform stakeholders, and helping scale the India AI engineering hub during the Wipro engagement through **300+ résumé screens** and **50+ technical interviews**.