

AI EVALUATION

- LLM Evaluation
- Prompt Configuration
- Accuracy
- Confusion Matrix
- MSE
- Structured JSON
- Multi-Task Testing

ENGINEERING REVIEW

- System Architecture
- Technical Design
- Impact Analysis
- Code Review
- Unit Testing
- Integration Testing
- Load Testing
- Debugging

SOFTWARE & AI

- Python
- Java
- SQL
- PyTorch
- HuggingFace
- PEFT
- Unsloth
- vLLM
- XGBoost
- Apache Beam

DELIVERY & QUALITY

- Git
- GitOps
- CI/CD
- Docker
- Kubernetes
- Prometheus
- Grafana
- Technical Documentation
- Knowledge Transfer

CERTIFICATIONS

- Oracle Cloud Infrastructure 2024
Generative AI Certified Professional
- Microsoft Certified
Azure AI Engineer Associate
- Microsoft Certified
Azure Data Scientist Associate
- Google Cloud
Certified Generative AI Leader

EDUCATION

B.E. Electrical & Electronics
Andhra University | 2012 - 2016

ML METHODS

- Scikit-learn
- TensorFlow
- Keras
- DenseNet
- Classification
- Predictive Modelling
- Unsupervised Learning
- Feature Engineering
- Transfer Learning
- Anomaly Detection
- NLP
- Time-Series Forecasting

Production AI and software engineer with 10+ years designing, reviewing, testing, and operating complex systems. At Best Buy, designed a config-driven LoRA framework with prompt files and automated evaluation using accuracy, confusion matrices, MSE, and structured JSON, then connected it to production vLLM serving. Lead AI/ML design, technical review, quality checks, debugging, and cross-team decisions for systems with feature pipelines handling **40,000 events/sec** at peak and a **200 ms scoring SLA**. Brings practical judgment from LLM, ML platform, cybersecurity, cloud, and client-delivery work.

EXPERIENCE

Lead Artificial Intelligence Engineer

Best Buy India

Jul 2025 - Present | Bengaluru

- > Designed a config-driven LoRA fine-tuning framework where new tasks are added through YAML, prompt files, and CSV datasets; supports Unsloth and HuggingFace PEFT with single- and multi-GPU execution.
- > Built automated post-training evaluation using accuracy, confusion matrices, MSE, and structured JSON results, with repeatable artifact handling and vLLM serve-command generation.
- > Architected and reviewed the production path from model training through shared vLLM serving on GKE, including HPA autoscaling, Istio routing, blue-green delivery, Prometheus, and Grafana.
- > Built fraud systems across batch, streaming, and real-time models. Beam/Dataflow feature pipelines handled **40,000 events/sec** at peak; supported a **200 ms scoring SLA** with secure, monitored GKE/Seldon Core inference and load testing.

Lead ML Engineer — Machine Learning & MLOps

Wipro

Aug 2023 - Jul 2025 | Hyderabad

- > Led design, review, debugging, and production delivery for cybersecurity ML across GCP, AWS workloads, and an NVIDIA DGX with 8xA100 GPUs. Completed **300+ résumé screens** and **50+ technical interviews** at Wipro.
- > Improved phishing-model throughput to **2,400 emails/sec**, twice the prior rate, and operated continuously refreshed behavioral models over billion-row datasets.
- > Reviewed system behavior and model outputs with engineering and client teams, mentored developers, and turned production issues into clear technical changes.

Senior Data Scientist — Data Science / Azure ML

Accenture

May 2021 - Sep 2023 | India

- > Translated client requirements into Azure ML designs, automated retraining, deployment workflows, technical demonstrations, and quality checks for supply-chain systems.
- > Led a product-quality system that increased inspections from **130 to 160 trucks/day**, improved low-quality detection 17%, and produced **\$2.5M annual savings**.

Earlier Software, ML & Data Engineering Roles

GlobalData Research | TCS | Infosys

May 2016 - May 2021 | India and South Korea

- > Delivered forecasting across 15+ industries, a distributed multi-GPU vision system trained on about 1M images, and mainframe software work spanning impact analysis, design, coding, review, unit tests, integration tests, and production support.

PROJECTS

LLM Evaluation Pipeline

Best Buy | Prompts | Metrics | Structured Results

Repeatable evaluation for classification, extraction, and generation tasks, with task-specific prompt files and machine-readable output for comparing model behavior.

Architecture Quality at Scale

Best Buy | 40k feature events/sec | 200 ms scoring SLA

Reviewed system design, load behavior, deployment safety, observability, and security for low-latency production ML serving.

GPU STACK

- CUDA GPU Libraries
- RAPIDS cuDF / CuPy
- Multi-GPU Training

OPTIMISATION & SIMULATION

- Google OR-Tools
- Vehicle Routing
- Mixed-Integer Programming
- Constraint Modelling
- Operations Research
- Assignment Models
- Supply-Chain Optimization
- Scenario Analysis

CLOUD & MANAGED ML

- GCS
- Artifact Registry
- Azure ML
- Azure DevOps
- Function Apps
- Logic Apps
- Cloud Build
- Google KCC

DATA & ML LIFECYCLE

- Dataflow
- Pub/Sub
- Vertex AI Feature Store
- Streaming Pipelines
- Training Pipelines
- Automated Evaluation
- Experiment Tracking
- Model Registry
- Model Versioning
- Continuous Training

LANGUAGES & DATA

- Pandas
- NumPy
- MS SQL Server
- IBM DB2
- Power BI

ENGINEERING LEADERSHIP

- Architecture Reviews
- Technical Mentoring
- Technical Hiring
- Team Building
- Stakeholder Alignment
- Cross-Functional Delivery

ACHIEVEMENTS

15 honors and awards across Best Buy, Wipro, Accenture, and Infosys.
Best Buy Talent Quotient (2026): cybersecurity and fraud delivery.
Best Buy client AML Hall of Fame (2024): DGX delivery during the Wipro engagement.

ADDITIONAL PROJECTS & EXPERIENCE

Model Error & Performance Work
Wipro | Cybersecurity | Billion-row datasets

Analyzed model and system behavior, improved throughput, and supported continuously refreshed ML across high-volume security workloads.

Technical Communication
Clients | Teams | Hiring

Explained designs through demos and knowledge transfer, mentored engineers, reviewed technical work, and conducted structured technical interviews.

EARLIER EXPERIENCE & ACHIEVEMENTS

Data Scientist
GD Research Centre | May 2020 - May 2021

Built domain-aware time-series forecasting and scenario analysis across energy, manufacturing, healthcare, finance, retail, telecommunications, automotive, and other industry intelligence centers.

Deep Learning Engineer
Tata Consultancy Services | Jul 2019 - Mar 2020

Built and deployed an end-to-end semiconductor ML system: distributed training on ~1 million images in ~3 hours, with 90% classification accuracy. Deployed at Amkor, South Korea, and trained the client team.

Senior Software Engineer - ML
Infosys | May 2016 - Jul 2019

Built an ML-powered lead-conversion prediction system and business-intelligence applications that moved sales prioritisation from intuition to data.

Achievements
15 honors and awards

Best Buy Talent Quotient (2026); Best Buy client AML Hall of Fame (2024), during the Wipro engagement.

Vehicle Routing & Transport Planning
Accenture | Google OR-Tools | Operations research

Designed vehicle-routing and mixed-integer models covering capacity, time windows, waiting time, demand, and departure constraints, then automated daily operational planning.

Industry Forecasting Platform
GD Research Centre | Time series | Scenario analysis

Designed and deployed domain-aware forecasting across 15+ industries, combining feature engineering, time-series modelling, and scenario analysis for enterprise decisions.

System Ownership
Design | Implementation | Production operations | Maintenance

Owned AI/ML systems from data and training pipelines through evaluation, deployment, monitoring, maintenance, incident support, and continuous improvement.

Engineering Leadership & Team Building
Architecture | Mentoring | Hiring | Cross-functional delivery

Stayed hands-on while leading ML/MLOps delivery, mentoring engineers, coordinating client and platform stakeholders, and helping scale the India AI engineering hub during the Wipro engagement through 300+ résumé screens and 50+ technical interviews.