

LAXMI NARASIMHA

Hari Yelesetty

Lead AI/ML Engineer & Data Scientist | Production ML & GenAI | Distributed Systems |
GCP · Azure · GPU Computing

ynharimailme@gmail.com
+91 99161 12881
linkedin.com/in/ynhari
github.com/ynhari
ynhari.github.io

ENGINEERING

Python SQL Distributed Systems
API Services Async Processing

APPLIED AI

vLLM LoRA / QLoRA PEFT
Transformers PyTorch Model Evaluation
Distributed Training Inference APIs

PRODUCTION DELIVERY

GCP Azure GKE Kubernetes Docker
Helm Istio Cloud Build ArgoCD

QUALITY & TRUST

GitOps CI/CD Prometheus Grafana
MLflow Load Testing
Blue-Green Delivery JWT Security

CERTIFICATIONS

Oracle Cloud Infrastructure 2024
Generative AI Certified Professional
Microsoft Certified
Azure AI Engineer Associate
Microsoft Certified
Azure Data Scientist Associate
Google Cloud
Certified Generative AI Leader

EDUCATION

B.E. Electrical & Electronics
Andhra University | 2012 - 2016

ML METHODS

Scikit-learn TensorFlow Keras
DenseNet Classification
Predictive Modelling
Unsupervised Learning
Feature Engineering Transfer Learning
Anomaly Detection NLP
Time-Series Forecasting

Lead AI and ML engineer and data scientist with **10+ years** of experience, building production AI/ML/LLM systems. Leads end-to-end AI delivery, combining hands-on model development with technical leadership to turn business problems into scalable, high-throughput, low-latency production systems. Experience spans predictive modelling, forecasting, operations research, computer vision, and generative AI across retail, cybersecurity, manufacturing, insurance, and supply-chain optimization.

EXPERIENCE

Lead Artificial Intelligence Engineer
Best Buy India

Jul 2025 - Present | Bengaluru

- Built a shared production LLM platform covering config-driven LoRA/QLoRA training, automated evaluation, artifact lifecycle, model caching, and vLLM OpenAI-compatible serving on GKE.
- Designed HPA autoscaling, Istio routing, blue-green releases, Cloud Build, ArgoCD, KCC, Prometheus, and Grafana; balanced quality, latency, reliability, governance, and cost with estimated annual API-cost savings of **~\$300K**. Kept model inputs company-managed and reduced external API exposure during peak demand.
- Built fraud systems across batch, streaming, and real-time models. Beam/Dataflow feature pipelines handled **40,000 events/sec** at peak; supported a **200 ms scoring SLA** with secure, monitored GKE/Seldon Core inference and load testing.
- Lead technical direction across data science, platform, fraud, and security teams; review designs and mentor engineers.

Lead ML Engineer — Machine Learning & MLOps
Wipro

Aug 2023 - Jul 2025 | Hyderabad

- Built distributed training and inference pipelines on NVIDIA DGX (8xA100), processing **billion-row** AWS CloudTrail datasets with Morpheus, RAPIDS, continuous retraining, and Grafana alerting.
- Engineered phishing classification at **2,400 emails/sec** and automated multiformat abuse-inbox workflows across Outlook APIs, OCR, analysts, and external response partners.
- Led and mentored ML/MLOps engineers; coordinated client/NVIDIA delivery; completed **300+ résumé screens** and conducted **50+ technical interviews**.

Senior Data Scientist — Data Science / Azure ML
Accenture

May 2021 - Sep 2023 | India

- Led end-to-end Azure ML delivery for BamaGruppen, increasing inspections from **130 to 160 trucks/day**, improving low-quality detection 17%, and generating **\$2.5M annual savings**.
- Translated logistics needs into operated Python services using Azure Functions, Logic Apps, Azure DevOps, CI/CD, and Google OR-Tools; demonstrated solutions, handled support, and tracked outcomes with client stakeholders.

PROJECTS

LLM Lifecycle Platform
Best Buy | Evaluation | GKE | vLLM

Shared Python platform for fine-tuning, automated evaluation, artifact management, model caching, autoscaled inference, safe rollout, and full-stack observability.

Real-Time Fraud Platform
Best Buy | 40k feature events/sec | 200 ms scoring SLA

Streaming feature engineering and low-latency inference for high-traffic retail workflows, with secure endpoints, load tests, telemetry, alerts, and GitOps delivery.

GPU STACK

- CUDA GPU Libraries
- RAPIDS cuDF / CuPy
- Multi-GPU Training

OPTIMISATION & SIMULATION

- Google OR-Tools
- Vehicle Routing
- Mixed-Integer Programming
- Constraint Modelling
- Operations Research
- Assignment Models
- Supply-Chain Optimization
- Scenario Analysis

CLOUD & MANAGED ML

- GCS
- Artifact Registry
- Azure ML
- Azure DevOps
- Function Apps
- Logic Apps
- Google KCC

DATA & ML LIFECYCLE

- Apache Beam
- Dataflow
- Pub/Sub
- Vertex AI Feature Store
- Streaming Pipelines
- Training Pipelines
- Automated Evaluation
- Experiment Tracking
- Model Registry
- Model Versioning
- Continuous Training

LANGUAGES & DATA

- Pandas
- NumPy
- MS SQL Server
- IBM DB2
- Power BI

ENGINEERING LEADERSHIP

- Architecture Reviews
- Technical Mentoring
- Technical Hiring
- Team Building
- Stakeholder Alignment
- Cross-Functional Delivery

ACHIEVEMENTS

15 honors and awards across Best Buy, Wipro, Accenture, and Infosys.
Best Buy Talent Quotient (2026): cybersecurity and fraud delivery.
Best Buy client AML Hall of Fame (2024): DGX delivery during the Wipro engagement.

ADDITIONAL PROJECTS & EXPERIENCE

Security AI Automation

Wipro | APIs | OCR | Human feedback

High-throughput classification and multiformat email processing connected to analysts and response partners, with monitored pipelines and continuous model improvement.

Client ML Delivery

Accenture | Azure | Python | APIs

Moved optimisation and quality-prediction systems from discovery to production services, automated pipelines, stakeholder demos, support, and measurable operating gains.

EARLIER EXPERIENCE & ACHIEVEMENTS

Data Scientist

GD Research Centre | May 2020 - May 2021

Built domain-aware time-series forecasting and scenario analysis across energy, manufacturing, healthcare, finance, retail, telecommunications, automotive, and other industry intelligence centers.

Deep Learning Engineer

Tata Consultancy Services | Jul 2019 - Mar 2020

Built and deployed an end-to-end semiconductor ML system: distributed training on **~1 million images in ~3 hours**, with **90% classification accuracy**. Deployed at Amkor, South Korea, and trained the client team.

Senior Software Engineer - ML

Infosys | May 2016 - Jul 2019

Built an ML-powered lead-conversion prediction system and business-intelligence applications that moved sales prioritisation from intuition to data.

Achievements

15 honors and awards

Best Buy Talent Quotient (2026); Best Buy client AML Hall of Fame (2024), during the Wipro engagement.

Vehicle Routing & Transport Planning

Accenture | Google OR-Tools | Operations research

Designed vehicle-routing and mixed-integer models covering capacity, time windows, waiting time, demand, and departure constraints, then automated daily operational planning.

Industry Forecasting Platform

GD Research Centre | Time series | Scenario analysis

Designed and deployed domain-aware forecasting across 15+ industries, combining feature engineering, time-series modelling, and scenario analysis for enterprise decisions.

System Ownership

Design | Implementation | Production operations | Maintenance

Owned AI/ML systems from data and training pipelines through evaluation, deployment, monitoring, maintenance, incident support, and continuous improvement.

Engineering Leadership & Team Building

Architecture | Mentoring | Hiring | Cross-functional delivery

Stayed hands-on while leading ML/MLOps delivery, mentoring engineers, coordinating client and platform stakeholders, and helping scale the India AI engineering hub during the Wipro engagement through **300+ résumé screens** and **50+ technical interviews**.