

AI INFRASTRUCTURE

- GPU Systems
- Distributed Systems
- Python
- Kubernetes
- Linux / Shell
- Production ML

MODEL PLATFORMS

- vLLM
- PyTorch
- TensorFlow Serving
- LoRA / QLoRA
- MLflow
- Seldon Core
- Batch / Online Inference

CLOUD & RUNTIME

- GCP
- Azure ML
- GKE
- Docker
- Helm
- Istio
- RAPIDS

DELIVERY & RELIABILITY

- Cloud Build
- ArgoCD
- GitOps
- Blue-Green Delivery
- Prometheus
- Grafana
- JWT Security
- Client Enablement

CERTIFICATIONS

- Oracle Cloud Infrastructure 2024
Generative AI Certified Professional
- Microsoft Certified
Azure AI Engineer Associate
- Microsoft Certified
Azure Data Scientist Associate
- Google Cloud
Certified Generative AI Leader

EDUCATION

- B.E. Electrical & Electronics
Andhra University | 2012 - 2016

ML METHODS

- Scikit-learn
- TensorFlow
- Keras
- DenseNet
- Classification
- Predictive Modelling
- Unsupervised Learning
- Feature Engineering
- Transfer Learning
- Anomaly Detection
- NLP
- Time-Series Forecasting

Hands-on AI/ML infrastructure engineer with 10+ years designing, deploying, scaling, and operating production systems across cloud and on-prem GPU environments. At Best Buy, architected shared vLLM serving and LoRA/QLoRA training on GKE, with estimated annual API-cost savings of **~\$300K**, and leads real-time ML delivery with feature pipelines handling **40,000 events/sec** at peak and a **200 ms scoring SLA**. Earlier delivered continuously retrained security workloads on an 8xA100 DGX, Azure ML solutions saving \$2.5M annually, and a multi-GPU vision system deployed at a client site in South Korea. Strong in Python, Kubernetes, Linux-based ML deployment, GPU platforms, automation, monitoring, troubleshooting, mentoring, and customer-facing production delivery.

EXPERIENCE

Lead Artificial Intelligence Engineer

Best Buy India

Jul 2025 - Present | Bengaluru

- Architected a reusable production LLM platform for config-driven LoRA/QLoRA training, automated evaluation, model artifacts, and shared vLLM OpenAI-compatible serving on GKE.
- Turned latency, capacity, reliability, security, and cost goals into HPA autoscaling, Istio routing, blue-green releases, Cloud Build, ArgoCD, Helm, Prometheus, and Grafana; self-hosted inference supports estimated annual API-cost savings of **~\$300K**. Kept model inputs company-managed and reduced external API exposure during peak demand.
- Built fraud systems across batch, streaming, and real-time models. Beam/Dataflow feature pipelines handled **40,000 events/sec** at peak; supported a **200 ms scoring SLA** with secure, monitored GKE/Seldon Core inference and load testing.
- Lead technical direction across data science, platform, fraud, and security teams; review designs and mentor engineers.

Lead ML Engineer — Machine Learning & MLOps

Wipro

Aug 2023 - Jul 2025 | Hyderabad

- Designed and operated GPU-accelerated AI systems on NVIDIA DGX (8xA100, 1TB RAM), running continuously refreshed training and inference across **billion-row** AWS CloudTrail datasets with Morpheus, RAPIDS, MLflow, and Grafana.
- Engineered production phishing classification at **2,400 emails/sec** - 2x prior throughput - plus GPU-accelerated OCR and multiformat security pipelines integrated with Microsoft Sentinel.
- Led and mentored ML/MLOps engineers; coordinated client/NVIDIA delivery; completed **300+ résumé screens** and conducted **50+ technical interviews**.

Senior Data Scientist — Data Science / Azure ML

Accenture

May 2021 - Sep 2023 | India

- Led Azure ML product delivery from requirements and architecture through automated retraining, deployment, monitoring, client demonstrations, and support; increased inspections from **130 to 160 trucks/day** and generated **\$2.5M annual savings**.
- Built production Python services and automation with Azure Functions, Logic Apps, Azure DevOps, Blob Storage, and CI/CD; translated complex customer needs into operated cloud solutions.

PROJECTS

Enterprise LLM Platform

Best Buy | GKE | vLLM | GPU Inference

Training, evaluation, artifacts, shared model APIs, autoscaling, traffic control, secure rollout, cost optimization, and production observability.

GPU Security Platform

Wipro | DGX 8xA100 | 1B+ rows

Distributed training and inference, performance engineering, continuous refresh, monitoring, troubleshooting, and production integration.

GPU STACK

CUDA GPU Libraries

RAPIDS cuDF / CuPy

Multi-GPU Training

OPTIMISATION & SIMULATION

Google OR-Tools

Vehicle Routing

Mixed-Integer Programming

Constraint Modelling

Operations Research

Assignment Models

Supply-Chain Optimization

Scenario Analysis

CLOUD & MANAGED ML

GCS

Artifact Registry

Azure DevOps

Function Apps

Logic Apps

Google KCC

DATA & ML LIFECYCLE

Apache Beam

Dataflow

Pub/Sub

Vertex AI Feature Store

Streaming Pipelines

Training Pipelines

Automated Evaluation

Experiment Tracking

Model Registry

Model Versioning

Continuous Training

LANGUAGES & DATA

SQL

Pandas

NumPy

MS SQL Server

IBM DB2

Power BI

ENGINEERING LEADERSHIP

Architecture Reviews

Technical Mentoring

Technical Hiring

Team Building

Stakeholder Alignment

Cross-Functional Delivery

ACHIEVEMENTS

15 honors and awards across Best Buy, Wipro, Accenture, and Infosys.
Best Buy Talent Quotient (2026): cybersecurity and fraud delivery.
Best Buy client AML Hall of Fame (2024): DGX delivery during the Wipro engagement.

ADDITIONAL PROJECTS & EXPERIENCE

Secure Real-Time ML
Best Buy | 40k feature events/sec | 200 ms scoring SLA
Streaming features, Kubernetes inference, JWT-protected services, load testing, telemetry, alerting, and controlled deployment.

On-Site Vision Deployment
TCS / Amkor | South Korea | 1M images
Built and deployed an end-to-end semiconductor ML system: distributed training on ~1 million images in ~3 hours, with 90% classification accuracy. Deployed at Amkor, South Korea, and trained the client team.

EARLIER EXPERIENCE & ACHIEVEMENTS

Data Scientist
GD Research Centre | May 2020 - May 2021
Built domain-aware time-series forecasting and scenario analysis across energy, manufacturing, healthcare, finance, retail, telecommunications, automotive, and other industry intelligence centers.

Deep Learning Engineer
Tata Consultancy Services | Jul 2019 - Mar 2020
Built and deployed an end-to-end semiconductor ML system: distributed training on ~1 million images in ~3 hours, with 90% classification accuracy. Deployed at Amkor, South Korea, and trained the client team.

Senior Software Engineer - ML
Infosys | May 2016 - Jul 2019
Built an ML-powered lead-conversion prediction system and business-intelligence applications that moved sales prioritisation from intuition to data.

Achievements
15 honors and awards
Best Buy Talent Quotient (2026); Best Buy client AML Hall of Fame (2024), during the Wipro engagement.

Vehicle Routing & Transport Planning
Accenture | Google OR-Tools | Operations research
Designed vehicle-routing and mixed-integer models covering capacity, time windows, waiting time, demand, and departure constraints, then automated daily operational planning.

Industry Forecasting Platform
GD Research Centre | Time series | Scenario analysis
Designed and deployed domain-aware forecasting across 15+ industries, combining feature engineering, time-series modelling, and scenario analysis for enterprise decisions.

System Ownership
Design | Implementation | Production operations | Maintenance
Owned AI/ML systems from data and training pipelines through evaluation, deployment, monitoring, maintenance, incident support, and continuous improvement.

Engineering Leadership & Team Building
Architecture | Mentoring | Hiring | Cross-functional delivery
Stayed hands-on while leading ML/MLOps delivery, mentoring engineers, coordinating client and platform stakeholders, and helping scale the India AI engineering hub during the Wipro engagement through 300+ résumé screens and 50+ technical interviews.