

MODEL PLATFORM

vLLMOpenAI-Compatible APIKGKE

KubernetesHPA IstioModel Caching

Adapter Hot-Loading

BACKEND & SYSTEMS

PythonAPI ServicesDistributed Systems

Apache BeamDataflowPub/Sub

RELIABILITY & DELIVERY

PrometheusGrafanaBlue-Green Delivery

Cloud BuildArgoCDGoogle KCC

GitOpsCI/CDLoad Testing

ML LIFECYCLE

LoRA / QLoRAPEFTTransformers

Automated EvaluationVertex AI

Seldon Core

CERTIFICATIONS

Oracle Cloud Infrastructure 2024
Generative AI Certified Professional

Microsoft Certified
Azure AI Engineer Associate

Microsoft Certified
Azure Data Scientist Associate

Google Cloud
Certified Generative AI Leader

EDUCATION

B.E. Electrical & Electronics

Andhra University | 2012 - 2016

ML METHODS

Scikit-learnPyTorchTensorFlow

KerasDenseNetClassification

Predictive Modelling

Unsupervised Learning

Feature EngineeringTransfer Learning

Anomaly DetectionNLP

Time-Series Forecasting

Production AI and distributed-systems engineer with 10+ years building dependable platforms under real traffic. At Best Buy, architected a shared model platform spanning training, evaluation, artifact management, adapter hot-loading, and vLLM serving on GKE, with estimated annual API-cost savings of ~\$300K. Also leads real-time ML infrastructure with feature pipelines handling 40,000 events/sec at peak and a 200 ms scoring SLA. Earlier built billion-row GPU pipelines on 8xA100 DGX. Hands-on with Python, Go, Kubernetes, developer-facing APIs, rollout safety, observability, and cross-team technical leadership.

EXPERIENCE

Lead Artificial Intelligence Engineer

Best Buy India
Jul 2025 - Present | Bengaluru

- > Architected a shared production platform for config-driven LoRA/QLoRA training, automated evaluation, artifact lifecycle, and vLLM OpenAI-compatible serving on GKE.
- > Created a zero-Python task-extension path using YAML, prompt files, and datasets; PEFT adapters hot-load into a shared base model instead of requiring separate deployments.
- > Built reliable delivery with HPA, Istio, blue-green releases, Cloud Build, ArgoCD, Google KCC, Prometheus, and Grafana; self-hosted services support estimated annual API-cost savings of ~\$300K. Kept model inputs company-managed and reduced external API exposure during peak demand.
- > Built fraud systems across batch, streaming, and real-time models. Beam/Dataflow feature pipelines handled 40,000 events/sec at peak; supported a 200 ms scoring SLA with secure, monitored GKE/Seldon Core inference and load testing.

Lead ML Engineer — Machine Learning & MLOps

Wipro
Aug 2023 - Jul 2025 | Hyderabad

- > Built continuously retrained distributed pipelines on NVIDIA DGX (8xA100), processing billion-row AWS CloudTrail datasets with Morpheus, RAPIDS, model monitoring, and Grafana alerting.
- > Engineered phishing classification at 2,400 emails/sec - 2x prior throughput - and multifformat abuse-inbox processing integrated with Microsoft Sentinel and production security workflows.
- > Led and mentored ML/MLOps engineers; coordinated client/NVIDIA delivery; completed 300+ résumé screens and conducted 50+ technical interviews.

Senior Data Scientist — Data Science / Azure ML

Accenture
May 2021 - Sep 2023 | India

- > Led end-to-end Azure ML product delivery for BamaGruppen, increasing inspections from 130 to 160 trucks/day, improving low-quality detection 17%, and generating \$2.5M annual savings.
- > Built production APIs and automation with Python, Azure Functions, Logic Apps, Azure DevOps, CI/CD, and Google OR-Tools; turned ambiguous client needs into operated systems.

PROJECTS

Shared Model Platform

Best Buy | GKE | vLLM | Developer experience

One interface for training, evaluation, artifact management, model caching, adapter hot-loading, autoscaled inference, traffic routing, safe rollout, and observability.

Real-Time Retail ML Platform

Best Buy | 40k feature events/sec | 200 ms scoring SLA

Streaming feature engineering and low-latency inference across high-traffic retail workflows, with autoscaling, load testing, secure endpoints, telemetry, and GitOps delivery.

GPU STACK

- CUDA GPU Libraries
- RAPIDS cuDF / CuPy
- Multi-GPU Training

OPTIMISATION & SIMULATION

- Google OR-Tools
- Vehicle Routing
- Mixed-Integer Programming
- Constraint Modelling
- Operations Research
- Assignment Models
- Supply-Chain Optimization
- Scenario Analysis

CLOUD & MANAGED ML

- GCS
- Artifact Registry
- Azure ML
- Azure DevOps
- Function Apps
- Logic Apps

DATA & ML LIFECYCLE

- Vertex AI Feature Store
- Streaming Pipelines
- Training Pipelines
- Experiment Tracking
- Model Registry
- Model Versioning
- Continuous Training

LANGUAGES & DATA

- SQL
- Pandas
- NumPy
- MS SQL Server
- IBM DB2
- Power BI

ENGINEERING LEADERSHIP

- Architecture Reviews
- Technical Mentoring
- Technical Hiring
- Team Building
- Stakeholder Alignment
- Cross-Functional Delivery

ACHIEVEMENTS

15 honors and awards across Best Buy, Wipro, Accenture, and Infosys.
Best Buy Talent Quotient (2026): cybersecurity and fraud delivery.
Best Buy client AML Hall of Fame (2024): DGX delivery during the Wipro engagement.

ADDITIONAL PROJECTS & EXPERIENCE

GPU Cybersecurity Pipelines

Wipro | DGX | Billion-row scale

Continuously refreshed behavioral models and high-throughput classification pipelines with GPU acceleration, monitoring, and analyst-facing production integrations.

Rover - Go Systems Project

Open source | Go | Real-time SSE

Built a zero-dependency remote command executor and project launcher in Go with persistent sessions and real-time server-sent event streaming.

EARLIER EXPERIENCE & ACHIEVEMENTS

Data Scientist

GD Research Centre | May 2020 - May 2021

Built domain-aware time-series forecasting and scenario analysis across energy, manufacturing, healthcare, finance, retail, telecommunications, automotive, and other industry intelligence centers.

Deep Learning Engineer

Tata Consultancy Services | Jul 2019 - Mar 2020

Built and deployed an end-to-end semiconductor ML system: distributed training on ~1 million images in ~3 hours, with 90% classification accuracy. Deployed at Amkor, South Korea, and trained the client team.

Senior Software Engineer - ML

Infosys | May 2016 - Jul 2019

Built an ML-powered lead-conversion prediction system and business-intelligence applications that moved sales prioritisation from intuition to data.

Achievements

15 honors and awards

Best Buy Talent Quotient (2026); Best Buy client AML Hall of Fame (2024), during the Wipro engagement.

Vehicle Routing & Transport Planning

Accenture | Google OR-Tools | Operations research

Designed vehicle-routing and mixed-integer models covering capacity, time windows, waiting time, demand, and departure constraints, then automated daily operational planning.

Industry Forecasting Platform

GD Research Centre | Time series | Scenario analysis

Designed and deployed domain-aware forecasting across 15+ industries, combining feature engineering, time-series modelling, and scenario analysis for enterprise decisions.

System Ownership

Design | Implementation | Production operations | Maintenance

Owned AI/ML systems from data and training pipelines through evaluation, deployment, monitoring, maintenance, incident support, and continuous improvement.

Engineering Leadership & Team Building

Architecture | Mentoring | Hiring | Cross-functional delivery

Stayed hands-on while leading ML/MLOps delivery, mentoring engineers, coordinating client and platform stakeholders, and helping scale the India AI engineering hub during the Wipro engagement through 300+ résumé screens and 50+ technical interviews.