

INFERENCE & PERFORMANCE

PythonPyTorchTensorFlowvLLM

Model ServingQuantization

KV-Cache FitLoad Testing

Latency / Throughput

Performance Analysis

CLOUD & DEPLOYMENT

GCPGKEVertex AIKubernetes

DockerHelmIstioSeldon Core

Azure MLAzure Functions

RELIABILITY & DELIVERY

HPACloud BuildArgoCDGoogle KCC

GitOpsCI/CDPrometheusGrafana

Istio RoutingBlue-Green Delivery

DISTRIBUTED & GPU SYSTEMS

Apache BeamDataflowPub/Sub

Real-Time StreamingRAPIDS

Distributed TrainingGPU Infrastructure

Model Evaluation

CERTIFICATIONS

Oracle Cloud Infrastructure 2024
Generative AI Certified Professional

Microsoft Certified
Azure AI Engineer Associate

Microsoft Certified
Azure Data Scientist Associate

Google Cloud
Certified Generative AI Leader

EDUCATION

B.E. Electrical & Electronics

Andhra University | 2012 - 2016

ML METHODS

Scikit-learnKerasDenseNet

ClassificationPredictive Modelling

Unsupervised Learning

Feature EngineeringTransfer Learning

Anomaly DetectionNLP

Time-Series Forecasting

Lead AI and ML engineer and data scientist with **10+ years** of experience, building production AI/ML/LLM systems. Leads end-to-end AI delivery, combining hands-on model development with technical leadership to turn business problems into scalable, high-throughput, low-latency production systems. Experience spans predictive modelling, forecasting, operations research, computer vision, and generative AI across retail, cybersecurity, manufacturing, insurance, and supply-chain optimization. Selected contributions include modelling on **billion-row security datasets**, supply-chain optimization generating **~\$2.5M** in savings, and self-hosted LLMs saving an estimated **~\$300K** annually while improving availability during peak demand.

EXPERIENCE

Lead Artificial Intelligence Engineer

Best Buy India

Jul 2025 - Present | Bengaluru

- > Architected a shared Python platform for LoRA/QLoRA training, automated evaluation, artifact lifecycle, and vLLM OpenAI-compatible serving on GKE; adapters hot-load into a shared base-model deployment.
- > Designed and analysed autoscaled inference delivery with HPA, Istio routing, blue-green releases, Cloud Build, ArgoCD, KCC, Prometheus, and Grafana; self-hosting supports estimated annual API-cost savings of **~\$300K**. Kept model inputs company-managed and reduced external API exposure during peak demand.
- > Built fraud systems across batch, streaming, and real-time models. Beam/Dataflow feature pipelines handled **40,000 events/sec** at peak; supported a **200 ms scoring SLA** with secure, monitored GKE/Seldon Core inference and load testing.
- > Translate workload, latency, reliability, and cost findings into architecture decisions across data science, platform, fraud, and security teams; review designs and mentor engineers.

Lead ML Engineer — Machine Learning & MLOps

Wipro

Aug 2023 - Jul 2025 | Hyderabad

- > Built distributed training and inference pipelines on NVIDIA DGX (8xA100), processing **billion-row** AWS CloudTrail datasets with Morpheus, RAPIDS, continuous retraining, and Grafana alerting.
- > Engineered phishing classification at **2,400 emails/sec** - 2x prior throughput - and multiformat abuse-inbox processing integrated with Microsoft Sentinel and production security workflows.
- > Led and mentored ML/MLOps engineers; coordinated client/NVIDIA delivery; completed **300+ résumé screens** and conducted **50+ technical interviews**.

Senior Data Scientist — Data Science / Azure ML

Accenture

May 2021 - Sep 2023 | India

- > Led end-to-end Azure ML product delivery for BamaGruppen, increasing inspections from **130 to 160 trucks/day**, improving low-quality detection 17%, and generating **\$2.5M annual savings**.
- > Built production services and automation with Python, Azure Functions, Logic Apps, Azure DevOps, CI/CD, and Google OR-Tools; translated ambiguous client needs into operated systems.

PROJECTS

vLLM Model Chooser

Open source | 115 models | A100-B200

Built a live GPU-to-model planner that filters vLLM-compatible models by VRAM, quantization, and KV-cache fit across accelerator classes.

LLM Training & Serving Platform

Best Buy | GKE | vLLM | Autoscaling

Shared Python platform for training, evaluation, artifact management, model caching, adapter-aware serving, traffic routing, safe rollout, and full-stack observability.

GPU STACK

- CUDA GPU Libraries
- RAPIDS cuDF / CuPy
- Multi-GPU Training

OPTIMISATION & SIMULATION

- Google OR-Tools
- Vehicle Routing
- Mixed-Integer Programming
- Constraint Modelling
- Operations Research
- Assignment Models
- Supply-Chain Optimization
- Scenario Analysis

CLOUD & MANAGED ML

- GCS
- Artifact Registry
- Azure DevOps
- Function Apps
- Logic Apps

DATA & ML LIFECYCLE

- Vertex AI Feature Store
- Streaming Pipelines
- Training Pipelines
- Automated Evaluation
- Experiment Tracking
- Model Registry
- Model Versioning
- Continuous Training

LANGUAGES & DATA

- SQL
- Pandas
- NumPy
- MS SQL Server
- IBM DB2
- Power BI

ENGINEERING LEADERSHIP

- Architecture Reviews
- Technical Mentoring
- Technical Hiring
- Team Building
- Stakeholder Alignment
- Cross-Functional Delivery

ACHIEVEMENTS

15 honors and awards across Best Buy, Wipro, Accenture, and Infosys.
Best Buy Talent Quotient (2026): cybersecurity and fraud delivery.
Best Buy client AML Hall of Fame (2024): DGX delivery during the Wipro engagement.

ADDITIONAL PROJECTS & EXPERIENCE

Real-Time Retail ML Platform
Best Buy | 40k feature events/sec | 200 ms scoring SLA

Streaming feature engineering and low-latency inference with autoscaling, load testing, secure endpoints, telemetry, and GitOps delivery.

Cybersecurity GPU Pipelines
Wipro | DGX | Billion-row scale

Continuously refreshed behavioral models and high-throughput classification pipelines with distributed GPU acceleration, monitoring, and production integrations.

EARLIER EXPERIENCE & ACHIEVEMENTS

Data Scientist
GD Research Centre | May 2020 - May 2021

Built domain-aware time-series forecasting and scenario analysis across energy, manufacturing, healthcare, finance, retail, telecommunications, automotive, and other industry intelligence centers.

Deep Learning Engineer
Tata Consultancy Services | Jul 2019 - Mar 2020

Built and deployed an end-to-end semiconductor ML system: distributed training on ~1 million images in ~3 hours, with 90% classification accuracy. Deployed at Amkor, South Korea, and trained the client team.

Senior Software Engineer - ML
Infosys | May 2016 - Jul 2019

Built an ML-powered lead-conversion prediction system and business-intelligence applications that moved sales prioritisation from intuition to data.

Achievements
15 honors and awards

Best Buy Talent Quotient (2026); Best Buy client AML Hall of Fame (2024), during the Wipro engagement.

Vehicle Routing & Transport Planning
Accenture | Google OR-Tools | Operations research

Designed vehicle-routing and mixed-integer models covering capacity, time windows, waiting time, demand, and departure constraints, then automated daily operational planning.

Industry Forecasting Platform
GD Research Centre | Time series | Scenario analysis

Designed and deployed domain-aware forecasting across 15+ industries, combining feature engineering, time-series modelling, and scenario analysis for enterprise decisions.

System Ownership
Design | Implementation | Production operations | Maintenance

Owned AI/ML systems from data and training pipelines through evaluation, deployment, monitoring, maintenance, incident support, and continuous improvement.

Engineering Leadership & Team Building
Architecture | Mentoring | Hiring | Cross-functional delivery

Stayed hands-on while leading ML/MLOps delivery, mentoring engineers, coordinating client and platform stakeholders, and helping scale the India AI engineering hub during the Wipro engagement through 300+ résumé screens and 50+ technical interviews.