

SOLUTION ARCHITECTURE

- AI / ML Systems
- Distributed Systems
- REST APIs
- Data Pipelines
- Python

AI DEPLOYMENT

- vLLM
- Model Serving
- LoRA / QLoRA
- Quantization
- Automated Evaluation
- MLflow
- Batch / Online Inference

CLOUD & HYBRID RUNTIME

- GCP
- Azure ML
- GKE
- Kubernetes
- Docker
- Istio
- Seldon Core

DELIVERY & OPERATIONS

- Client Demonstrations
- Technical Leadership
- CI/CD
- ArgoCD
- Blue-Green Delivery
- Prometheus
- Grafana
- JWT Security

CERTIFICATIONS

- Oracle Cloud Infrastructure 2024
Generative AI Certified Professional
- Microsoft Certified
Azure AI Engineer Associate
- Microsoft Certified
Azure Data Scientist Associate
- Google Cloud
Certified Generative AI Leader

EDUCATION

B.E. Electrical & Electronics
Andhra University | 2012 - 2016

ML METHODS

- Scikit-learn
- PyTorch
- TensorFlow
- Keras
- DenseNet
- Classification
- Predictive Modelling
- Unsupervised Learning
- Feature Engineering
- Transfer Learning
- Anomaly Detection
- NLP
- Time-Series Forecasting

AI platform and solution-architecture engineer with 10+ years taking ML systems from requirements and design through deployment, performance tuning, and production support across cloud and on-prem GPU infrastructure. At Best Buy, architected shared vLLM serving on GKE and LoRA/QLoRA training, with estimated annual API-cost savings of **~\$300K**, and leads real-time AI delivery with feature pipelines handling **40,000 events/sec** at peak and a **200 ms scoring SLA**. Earlier delivered billion-row pipelines on an 8xA100 DGX, client-facing Azure ML products saving \$2.5M annually, and an on-site semiconductor vision deployment in South Korea. Hands-on with Python, Kubernetes, GCP, Azure, distributed systems, model APIs, GPU performance, security, observability, technical reviews, mentoring, client demonstrations, and production handover.

EXPERIENCE

Lead Artificial Intelligence Engineer

Best Buy India
Jul 2025 - Present | Bengaluru

- Architected a reusable production LLM platform for config-driven LoRA/QLoRA training, automated evaluation, artifact lifecycle, and vLLM OpenAI-compatible serving on GKE.
- Turned latency, capacity, security, reliability, and cost needs into HPA autoscaling, Istio routing, blue-green releases, Cloud Build, ArgoCD, Prometheus, and Grafana; self-hosted inference supports estimated annual API-cost savings of **~\$300K**. Kept model inputs company-managed and reduced external API exposure during peak demand.
- Built fraud systems across batch, streaming, and real-time models. Beam/Dataflow feature pipelines handled **40,000 events/sec** at peak; supported a **200 ms scoring SLA** with secure, monitored GKE/Seldon Core inference and load testing.
- Lead technical direction across data science, platform, fraud, and security teams; review designs and mentor engineers.

Lead ML Engineer — Machine Learning & MLOps

Wipro
Aug 2023 - Jul 2025 | Hyderabad

- Designed hybrid GPU AI systems on NVIDIA DGX (8xA100), running continuously refreshed training and inference over **billion-row** AWS CloudTrail datasets with Morpheus, RAPIDS, MLflow, and Grafana.
- Engineered production phishing classification at **2,400 emails/sec** - 2x prior throughput - plus GPU-accelerated multifORMAT security workflows integrated with Microsoft Sentinel.
- Led and mentored ML/MLOps engineers; coordinated client/NVIDIA delivery; completed **300+ résumé screens** and conducted **50+ technical interviews**.

Senior Data Scientist — Data Science / Azure ML

Accenture
May 2021 - Sep 2023 | India

- Led Azure ML solution delivery from requirements and architecture through retraining, deployment, monitoring, client demonstrations, and support; increased inspections from **130 to 160 trucks/day** and generated **\$2.5M annual savings**.
- Built production Python services and automation with Azure Functions, Logic Apps, Azure DevOps, CI/CD, Blob Storage, and Google OR-Tools; converted complex client needs into operated cloud products.

PROJECTS

Enterprise LLM Platform

Best Buy | GKE | vLLM | GPU Inference

Training, evaluation, artifacts, shared model APIs, autoscaling, traffic management, controlled rollout, security, and observability across platform and product teams.

GPU-to-Model Planning

Open source | 115 models | A100-B200

Built a live planner that compares vLLM-compatible models by accelerator memory, quantization, and KV-cache fit across GPU classes.

GPU STACK

- CUDA GPU Libraries
- RAPIDS cuDF / CuPy
- Multi-GPU Training

OPTIMISATION & SIMULATION

- Google OR-Tools
- Vehicle Routing
- Mixed-Integer Programming
- Constraint Modelling
- Operations Research
- Assignment Models
- Supply-Chain Optimization
- Scenario Analysis

CLOUD & MANAGED ML

- GCS
- Artifact Registry
- Azure DevOps
- Function Apps
- Logic Apps
- Cloud Build
- Google KCC

DATA & ML LIFECYCLE

- Apache Beam
- Dataflow
- Pub/Sub
- Vertex AI Feature Store
- Streaming Pipelines
- Training Pipelines
- Experiment Tracking
- Model Registry
- Model Versioning
- Continuous Training

LANGUAGES & DATA

- SQL
- Pandas
- NumPy
- MS SQL Server
- IBM DB2
- Power BI

ENGINEERING LEADERSHIP

- Architecture Reviews
- Technical Mentoring
- Technical Hiring
- Team Building
- Stakeholder Alignment
- Cross-Functional Delivery

ACHIEVEMENTS

15 honors and awards across Best Buy, Wipro, Accenture, and Infosys.
Best Buy Talent Quotient (2026): cybersecurity and fraud delivery.
Best Buy client AML Hall of Fame (2024): DGX delivery during the Wipro engagement.

ADDITIONAL PROJECTS & EXPERIENCE

Hybrid GPU Security Platform
Wipro | DGX | Billion-row scale

Distributed training and inference, performance engineering, continuous refresh, monitoring, and production integration on on-prem GPU infrastructure.

On-Site AI Deployment
TCS / Amkor | South Korea | 1M images

Built and deployed an end-to-end semiconductor ML system: distributed training on **~1 million images in ~3 hours**, with **90% classification accuracy**. Deployed at Amkor, South Korea, and trained the client team.

EARLIER EXPERIENCE & ACHIEVEMENTS

Data Scientist
GD Research Centre | May 2020 - May 2021

Built domain-aware time-series forecasting and scenario analysis across energy, manufacturing, healthcare, finance, retail, telecommunications, automotive, and other industry intelligence centers.

Deep Learning Engineer
Tata Consultancy Services | Jul 2019 - Mar 2020

Built and deployed an end-to-end semiconductor ML system: distributed training on **~1 million images in ~3 hours**, with **90% classification accuracy**. Deployed at Amkor, South Korea, and trained the client team.

Senior Software Engineer - ML
Infosys | May 2016 - Jul 2019

Built an ML-powered lead-conversion prediction system and business-intelligence applications that moved sales prioritisation from intuition to data.

Achievements
15 honors and awards

Best Buy Talent Quotient (2026); Best Buy client AML Hall of Fame (2024), during the Wipro engagement.

Vehicle Routing & Transport Planning
Accenture | Google OR-Tools | Operations research

Designed vehicle-routing and mixed-integer models covering capacity, time windows, waiting time, demand, and departure constraints, then automated daily operational planning.

Industry Forecasting Platform
GD Research Centre | Time series | Scenario analysis

Designed and deployed domain-aware forecasting across 15+ industries, combining feature engineering, time-series modelling, and scenario analysis for enterprise decisions.

System Ownership
Design | Implementation | Production operations | Maintenance

Owned AI/ML systems from data and training pipelines through evaluation, deployment, monitoring, maintenance, incident support, and continuous improvement.

Engineering Leadership & Team Building
Architecture | Mentoring | Hiring | Cross-functional delivery

Stayed hands-on while leading ML/MLOps delivery, mentoring engineers, coordinating client and platform stakeholders, and helping scale the India AI engineering hub during the Wipro engagement through **300+ résumé screens** and **50+ technical interviews**.